

## Assignment 8 - SGD using linear Regression

```
In [1]: 1  ## Lets import Packages from the required libraries
        2  from sklearn.datasets import load_boston
        3  from sklearn.model_selection import train_test_split
        4  from sklearn.preprocessing import StandardScaler
        5  from sklearn.linear_model import SGDRegressor
        6  from sklearn.linear_model import LinearRegression
        7  from sklearn.metrics import mean_squared_error
        8  import pandas as pd
        9  import seaborn as sns
       10  import matplotlib.pyplot as plt
       11  import numpy as np
```

```
In [2]: 1  ## Loading our dataset
        2  boston=load_boston()
```

```
In [3]: 1  # Shape of the dataset
        2  boston.data.shape
```

Out[3]: (506, 13)

```
In [4]: 1  print(boston.feature_names)

['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
'B' 'LSTAT']
```

In [5]:

```
1 ## Description of each feature
2 print(boston.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.
```

```
:Attribute Information (in order):
```

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B  $1000(Bk - 0.63)^2$  where Bk is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

```
This is a copy of UCI ML housing dataset.
```

```
https://archive.ics.uci.edu/ml/machine-learning-databases/housing/ (https://archive.ics.uci.edu/ml/machine-learning-databases/housing/)
```

```
This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.
```

```
The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.
```

```
The Boston house-price data has been used in many machine learning papers that address regression problems.
```

```
.. topic:: References
```

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan,R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [6]:

```
1 boston_data=pd.DataFrame(data=boston.data,columns=boston.feature_names)
```

```
In [7]: 1 boston_data['price']=boston.target
        2 boston_data.head(2)
```

Out[7]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.9	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.9	9.14	21.6

Train Test Split

```
In [8]: 1 x=boston.data
        2 y=boston.target
        3 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
```

Standardize the Train set

```
In [9]: 1 ## using standard scalar starndardize the train data and tranform the test later to prevent data Leakage
        2 x=boston.data
        3 y=boston.target
        4
        5 scaler=StandardScaler()
        6
        7 scaler.fit(x_train)
        8 x_train=scaler.transform(x_train)
        9 x_test=scaler.transform(x_test)
       10 # Creating a new dataframe with tranformed data
       11 transformed_train_set=pd.DataFrame(x_train)
       12 transformed_train_set['price']=y_train
       13
```

Write custom SGD for linear Regression

```

In [10]: 1 class SGD_scratch :
2         '''The class contains various functions to perform Schocastic Gradient Descent on Linear regression
3         Class includes 2 functions primarily
4         1.find_w_b function : To find best (w,b) that fits the data giving minimum loss using SGD.
5         Note: Some important parameters from sklearn.SGDRegressor have also been implemented
6         2.Predict Function : that predicts y label given (w,b) '''
7     def __init__(self,dataset,learning_rate,k,no_iterations,shuffle=False,divide=None):
8         # initiating values
9         self.data=dataset
10        self.learning_rate=learning_rate
11        self.k=k
12        self.divide=divide
13        self.no_iterations=no_iterations
14        self.shuffle=shuffle
15
16    def find_w_b(self):
17        from sklearn.utils import shuffle
18        k,no_iter=self.k,self.no_iterations
19        w=np.random.randn(1,13)
20        b=0
21        lr_rate=self.learning_rate
22
23        for i in range(no_iter):
24            w_vector,b_vector=np.zeros(shape=(1,13)),0
25            if shuffle:
26                self.data=shuffle(self.data) #Shuffle the dataframe when shuffle=True for every iterations
27            else:
28                pass
29
30
31            #take sample of size k and store the data and target seprately
32            sampled_data=self.data.sample(k,replace=False)
33            x=np.array(sampled_data.drop('price',axis=1))
34            y=np.array(sampled_data['price'])
35            residual_sum_square=0
36            for pt in range(k):
37                actual_=y[pt] ## actual value y[i]
38                predicts_ = np.dot(w,x[pt]) + b ## predicted value w.T.x[pt] + intercept
39                w_vector+=(-2) * x[pt] * (y[pt] - (np.dot(w,x[pt]) + b) )
40                b_vector+=(-2) * (y[pt] - (np.dot(w,x[pt]) + b))
41                error = (predicts_ - actual_) ** 2 # square( actual(y[i]) - predicted(y[i]) )
42                residual_sum_square += error # sum of square of all the residue
43
44            print("-"*20)
45            print('Iteration {0} : mean squared error : {1} '.format(i,residual_sum_square/k))
46
47            previous_state_b=b # storing the previous intercept value
48            previous_state_w=w #storing the previous weights values
49
50            w=w-lr_rate*(w_vector/k)
51            b=b-lr_rate*(b_vector/k)
52
53            diff_vec= previous_state_w - w #calculate the difference between the vectors
54            if abs(previous_state_b -b) < 0.0008 and np.sqrt(np.sum(diff_vec**2)) < 0.0008 : # exit loop if w,b value is almost same,set random threshold 0.008
55                print('*'*20)
56                print('State wj :', previous_state_w)
57                print('State wj+1 :', w)
58                print('State bj :', previous_state_b)
59                print('State bj+1 :', b)
60                print('Iteration terminated at :',i)

```

```

61         print('Iteration terminated at :',i)
62         break
63
64     if self.divide:
65         lr_rate = lr_rate / pow((i+1),0.05)
66         print('lr_rate: ',lr_rate)
67     print("-"*10)
68
69     return w,b
70
71     def predict_func(self,w,b,test_set):
72         predict_set=[]
73         test_set=np.array(test_set)
74         for pt in range(len(test_set)):
75             predictions=(np.dot(w,test_set[pt])+b)
76             predict_set.append(predictions[0])
77         return predict_set

```


## Sklearn SGD Regressor

```

In [46]: 1 # SkLearn SGD classifier
2 n_iter=500
3 sklearn_model = SGDRegressor(max_iter=n_iter,alpha=0.05,shuffle=True)
4 sklearn_model.fit(x_train, y_train)
5 y_pred=sklearn_model.predict(x_test)
6 weights= pd.DataFrame(data=sklearn_model.coef_.T,columns=[ 'Sklearn SGD' ])
7 ## Plotting functions
8 #Scatter plot
9 predict_data=pd.DataFrame(data= {'Actual' : y_test,'Predicted' :y_pred})
10 sns.regplot(x='Actual', y='Predicted', data=predict_data)
11
12 print('Mean Squared Error :',mean_squared_error(y_test, y_pred))

```

Mean Squared Error : 22.2971947293




## Sklearn Linear Regression

```

In [47]: ▶ 1 # SkLearn SGD classifier
2 sklearn_lr_model = LinearRegression()
3 sklearn_lr_model.fit(x_train, y_train)
4 y_pred=sklearn_lr_model.predict(x_test)
5 weights['Sklearn Linear Regression '] = pd.DataFrame(sklearn_lr_model.coef_.T)
6 ## Plotting functions
7 #Scatter plot
8 predict_data=pd.DataFrame(data= {'Actual' : y_test,'Predicted' :y_pred})
9 sns.regplot(x='Actual', y='Predicted', data=predict_data)
10
11 print('Mean Squared Error :',mean_squared_error(y_test, y_pred))

```

Mean Squared Error : 22.1950284502



## Custom SGD regressor : Trial 1

1. Initially implementing with learning rate =1 and iterations=100
2. For every iteration set learning rate = learning rate /2
3. Shuffle the dataset per iteration

```
In [13]: 1 model=SGD_scratch(transformed_train_set,1,40,100,shuffle=True,divide=2)
2 w,b=model.find_w_b()
3 predictions=model.predict_func(w,b,x_test)
4 weights['Custom SGD (Trial 1)'] = pd.DataFrame(w.T)
```

```
lr_rate: 0.6529478154043212
```

```
-----
```

```
Iteration 7 : mean squared error : [ 8.57695860e+15]
```

```
lr_rate: 0.5884695206938755
```

```
-----
```

```
Iteration 8 : mean squared error : [ 2.98338967e+17]
```

```
lr_rate: 0.5272442454241162
```

```
-----
```

```
Iteration 9 : mean squared error : [ 1.02650061e+19]
```

```
lr_rate: 0.46990692835986236
```

```
-----
```

```
Iteration 10 : mean squared error : [ 4.48201807e+20]
```


```
lr_rate: 0.41681391978913523
```

```
-----
```

```
Iteration 11 : mean squared error : [ 9.33673670e+21]
```

```
In [14]: 1 ## Plotting functions
2 #Scatter plot
3 predict_data=pd.DataFrame(data= {'Actual' : y_test,'Predicted' :predictions})
4 sns.regplot(x='Actual', y='Predicted', data=predict_data)
5 print('Mean Squared Error :',mean_squared_error(y_test, predictions))
```

```
Mean Squared Error : 8.36279966769e+24
```



```
In [15]: 1 plt.figure(figsize=(25,6))
2 plt.plot(y_test, label='Actual')
3 plt.plot(predictions, label='Predicted')
4 plt.legend(prop={'size': 16})
5 plt.show()
6 print('Mean Squared Error :',mean_squared_error(y_test, predictions))
```



Mean Squared Error : 8.36279966769e+24

## Observations :

1. We see that the model does not perform well and gives a MSE of 8.36 e+24, with the above set parameters

## Custom SGD regressor :Trial 2

1. Implementing with learning rate =0.2 and iterations=500
2. Shuffle the dataset per iteration



```
In [55]: 1 model=SGD_scratch(transformed_train_set,0.2,40,500,shuffle=True,divide=1)
2 w,b=model.find_w_b()
3 predictions=model.predict_func(w,b,x_test)
4 weights['Custom SGD (Trial 2)'] = pd.DataFrame(w.T)
```

```
-----
Iteration 0 : mean squared error : [ 700.84819584]
lr_rate: 0.2
-----
```

```
-----
Iteration 1 : mean squared error : [ 316.38821338]
lr_rate: 0.19318726578496911
-----
```

```
-----
Iteration 2 : mean squared error : [ 441.89517938]
lr_rate: 0.18286156535236558
-----
```

```
-----
Iteration 3 : mean squared error : [ 629.41243007]
lr_rate: 0.17061587335782108
-----
```

```
-----
Iteration 4 : mean squared error : [ 1002.15283167]
lr_rate: 0.15742399642419647
-----
```

```
-----
Iteration 5 : mean squared error : [ 936.9914339]
lr_rate: 0.14393399205076884
-----
```

```
-----
Iteration 6 : mean squared error : [ 467.33421729]
lr_rate: 0.13058956308086425
-----
```

```
-----
Iteration 7 : mean squared error : [ 187.10058905]
lr_rate: 0.1176939041387751
-----
```

```
-----
Iteration 8 : mean squared error : [ 105.94217708]
lr_rate: 0.10544884908482326
-----
```

```
-----
Iteration 9 : mean squared error : [ 51.12823438]
lr_rate: 0.09398138567197249
-----
```

```
-----
Iteration 10 : mean squared error : [ 31.43511732]
lr_rate: 0.08336278395782706
-----
```

```
-----
Iteration 11 : mean squared error : [ 25.61902904]
lr_rate: 0.07362293851363488
-----
```

```
-----
Iteration 12 : mean squared error : [ 34.78255225]
lr_rate: 0.06476136289387929
-----
```

```
-----
Iteration 13 : mean squared error : [ 32.92859496]
lr_rate: 0.05675571223453064
```


```
-----  
-----  
Iteration 14 : mean squared error : [ 18.10594608]  
lr_rate: 0.04956841305788922  
-----  
-----  
Iteration 15 : mean squared error : [ 17.81742436]  
lr_rate: 0.04315180990924041  
-----  
-----  
Iteration 16 : mean squared error : [ 52.10837854]  
lr_rate: 0.03745213411452443  
-----  
-----  
Iteration 17 : mean squared error : [ 54.09439634]  
lr_rate: 0.03241253096307968  
-----  
-----  
Iteration 18 : mean squared error : [ 23.48152295]  
lr_rate: 0.02797533298812906  
-----  
-----  
Iteration 19 : mean squared error : [ 24.45075207]  
lr_rate: 0.024083730836508242  
-----  
-----  
Iteration 20 : mean squared error : [ 28.34336721]  
lr_rate: 0.020682965145182638  
-----  
-----  
Iteration 21 : mean squared error : [ 28.70609073]  
lr_rate: 0.017721140385663577  
-----  
-----  
Iteration 22 : mean squared error : [ 33.12278788]  
lr_rate: 0.015149743303671968  
-----  
-----  
Iteration 23 : mean squared error : [ 17.14495377]  
lr_rate: 0.012923933424513973  
-----  
-----  
Iteration 24 : mean squared error : [ 13.78995902]  
lr_rate: 0.011002660480289505  
-----  
-----  
Iteration 25 : mean squared error : [ 15.31844472]  
lr_rate: 0.009348653091118545  
-----  
-----  
Iteration 26 : mean squared error : [ 27.17433044]  
lr_rate: 0.007928314257640278  
-----  
-----  
Iteration 27 : mean squared error : [ 20.48470372]  
lr_rate: 0.006711551925354233  
-----  
-----  
Iteration 28 : mean squared error : [ 13.99175996]  
lr_rate: 0.005671566840000688  
-----
```

```
-----
Iteration 29 : mean squared error : [ 15.27902335]
lr_rate: 0.004784614938798966
-----
-----
Iteration 30 : mean squared error : [ 20.8187387]
lr_rate: 0.004029757453381585
-----
-----
Iteration 31 : mean squared error : [ 19.11246344]
lr_rate: 0.0033886085968905127
-----
-----
Iteration 32 : mean squared error : [ 13.30765294]
lr_rate: 0.002845088048680407
-----
-----
Iteration 33 : mean squared error : [ 30.13956293]
lr_rate: 0.002385183330304127
-----
-----
Iteration 34 : mean squared error : [ 22.7648264]
lr_rate: 0.001996725495132476
-----
-----
Iteration 35 : mean squared error : [ 33.75817345]
lr_rate: 0.00166918025298942
-----
-----
Iteration 36 : mean squared error : [ 19.68642721]
lr_rate: 0.001393455653778815
-----
-----
Iteration 37 : mean squared error : [ 40.08183626]
lr_rate: 0.001161726703474252
-----
-----
Iteration 38 : mean squared error : [ 22.5424867]
lr_rate: 0.0009672767341122829
-----
-----
Iteration 39 : mean squared error : [ 19.71794527]
lr_rate: 0.0008043549563845681
-----
-----
Iteration 40 : mean squared error : [ 19.30155866]
lr_rate: 0.0006680493556141013
-----
-----
Iteration 41 : mean squared error : [ 10.77947171]
lr_rate: 0.0005541739216873886
-----
-----
Iteration 42 : mean squared error : [ 41.22611897]
lr_rate: 0.0004591691082456614
-----
-----
Iteration 43 : mean squared error : [ 14.04456581]
lr_rate: 0.00038001437770191166
-----
-----
```


```
Iteration 44 : mean squared error : [ 27.34772574]
lr_rate: 0.00031415169161586796
-----
-----
Iteration 45 : mean squared error : [ 25.92401596]
lr_rate: 0.0002594188387963494
-----
-----
Iteration 46 : mean squared error : [ 24.32942131]
lr_rate: 0.00021399154687634468
-----
-----
Iteration 47 : mean squared error : [ 18.12495464]
lr_rate: 0.00017633338974736067
-----
-----
Iteration 48 : mean squared error : [ 33.92756921]
lr_rate: 0.0001451525775591791
-----
-----
Iteration 49 : mean squared error : [ 43.50212676]
lr_rate: 0.00011936479377123923
-----
-----
Iteration 50 : mean squared error : [ 13.43127502]
*****
State wj : [[-0.625174    1.08020076 -1.39007842  1.2845673  -0.40765815  2.93915697
 -0.50619766 -2.86436668  0.98714701 -0.70267729 -1.45645257  0.56144044
 -3.45457129]]
State wj+1 : [[-0.62514733  1.07999767 -1.3900024  1.28460618 -0.40767812  2.93881272
 -0.50603737 -2.86438383  0.98706158 -0.70279323 -1.45648774  0.56146142
 -3.45412598]]
State bj : [ 22.49808336]
State bj+1 : [ 22.49809334]
Iteration terminated at : 50
Iteration terminated at : 50
```

```
In [56]: 1 ## Plotting functions
2 #Scatter plot
3 predict_data=pd.DataFrame(data= {'Actual' : y_test,'Predicted' :predictions})
4 sns.regplot(x='Actual', y='Predicted', data=predict_data)
5 print('Mean Squared Error :',mean_squared_error(y_test, predictions))
```

Mean Squared Error : 22.5748669474



```
In [57]: 1 plt.figure(figsize=(25,6))
2 plt.plot(y_test, label='Actual')
3 plt.plot(predictions, label='Predicted')
4 plt.legend(prop={'size': 16})
5 plt.show()
6 print('Mean Squared Error :',mean_squared_error(y_test, predictions))
```



Mean Squared Error : 22.5748669474

## Observations :

1. We see that the model performs well and gives a MSE of 29.3 as close as sklearn's SGDRegressor, with the above set parameters

## Effect of learning rate on model performance

1. Considering a set of learning rate and check the model performance of each with constant iteration 100

```
In [40]: ▶ 1 alpha_set = [0.1,0.08,0.06,0.05,0.01]
2 custom_sgd_mse=[]
3 sk_sgd_mse=[]
4 for i in alpha_set:
5     # train custom model
6     model=SGD_scratch(transformed_train_set,i,40,500,shuffle=True,divide=1)
7     w,b=model.find_w_b()
8     predictions=model.predict_func(w,b,x_test)
9     custom_sgd_mse.append(mean_squared_error(y_test, predictions))
10
11     ##train SGDRegressor
12     sklearn_model = SGDRegressor(max_iter=500,alpha=i)
13     sklearn_model.fit(x_train, y_train)
14     y_pred=sklearn_model.predict(x_test)
15     sk_sgd_mse.append(mean_squared_error(y_test, y_pred))
```

```
lr_rate: 0.007574871651835984
```

```
-----
```

```
Iteration 23 : mean squared error : [ 19.99525915]
```

```
lr_rate: 0.0064619667122569864
```

```
-----
```

```
Iteration 24 : mean squared error : [ 29.64388956]
```

```
lr_rate: 0.005501330240144753
```

```
-----
```

```
Iteration 25 : mean squared error : [ 45.95801426]
```

```
lr_rate: 0.004674326545559273
```

```
-----
```

```
Iteration 26 : mean squared error : [ 9.25083364]
```


```
lr_rate: 0.003964157128820139
```

```
-----
```

```
-----
```

## Plotting MSE for custom SGD vs Sklearn SGD

```
In [41]: ▶ 1 plt.figure(figsize=(20,5))
2 plt.plot(alpha_set, custom_sgd_mse, color='blue', label='Custom SGD')
3 plt.plot(alpha_set, sk_sgd_mse, color='yellow', label='SK SGD')
4 plt.xlabel('Learning rates')
5 plt.ylabel('MSE')
6 plt.title("Comparison between the MSE of SKlearn SGD Regressor and Custom regressor for differnt Learning rates")
7 plt.legend()
8 plt.show()
```



## Observations

1. We can observe that our custom model has effect/fluctuations of learning rate to a greater extent than sklearn SGD.
2. With lower values of learning rate tends to decrease MSE (i.e learning rate  $< 0.06$  )

## Effect of number of iterations on model performance

1. Considering a set of iterations rate and check the model performance of each with constant learning rate 0.05

```
In [42]: 1 iter_set = [100,200,500,600,1000]
2 custom_sgd_mse=[]
3 sk_sgd_mse=[]
4 for i in iter_set:
5     # train custom model
6     model=SGD_scratch(transformed_train_set,0.05,40,i,shuffle=True,divide=1)
7     w,b=model.find_w_b()
8     predictions=model.predict_func(w,b,x_test)
9     custom_sgd_mse.append(mean_squared_error(y_test, predictions))
10
11     ##train SGDRegressor
12     sklearn_model = SGDRegressor(max_iter=i,alpha=0.05,shuffle=True)
13     sklearn_model.fit(x_train, y_train)
14     y_pred=sklearn_model.predict(x_test)
15     sk_sgd_mse.append(mean_squared_error(y_test, y_pred))
```

Iteration 5 : mean squared error : [ 235.62347647]

lr\_rate: 0.03598349801269221

-----

Iteration 6 : mean squared error : [ 195.01186375]

lr\_rate: 0.03264739077021606

-----

Iteration 7 : mean squared error : [ 163.57443764]

lr\_rate: 0.029423476034693776

-----

Iteration 8 : mean squared error : [ 149.84551938]

lr\_rate: 0.026362212271205814

-----

Iteration 9 : mean squared error : [ 125.73590188]

lr\_rate: 0.023495346417993123


-----

-----

## Plotting MSE for custom SGD vs Sklearn SGD



```
In [43]: ▶ 1 plt.figure(figsize=(20,5))
2 plt.plot(iter_set, custom_sgd_mse, color='blue', label='Custom SGD')
3 plt.plot(iter_set, sk_sgd_mse, color='yellow', label='SK SGD')
4 plt.xlabel('Iterations')
5 plt.ylabel('MSE')
6 plt.title("Comparison between the MSE of SKlearn SGD Regressor and Custom regressor for differnt Iterations")
7 plt.legend()
8 plt.show()
```



## Observations

1. We can observe that our custom model has effect/fluctuations of iterations to a greater extent than sklearn SGD.
2. As number of iterations increase it converges to optimum better.

## Comparing Weights and MSE

In [44]: ▶

1 weights

Out[44]:

	Sklearn SGD	Sklearn Linear Regression	Custom SGD (Trial 1)	Custom SGD (Trial 2)
0	-0.712181	-0.863947	1.652214e+12	0.088773
1	0.548530	0.743271	-1.479200e+12	0.999658
2	-0.709605	-0.412015	-1.618272e+12	-0.749712
3	1.082881	1.000662	-2.599226e+11	0.912983
4	-1.573638	-2.252131	-9.759675e+11	-0.076970
5	2.437577	2.217033	-3.254227e+10	3.052044
6	-0.056164	0.065871	-9.869387e+11	0.302897
7	-2.604154	-3.255986	-7.281176e+11	-1.504199
8	1.513546	2.875749	6.985335e+11	-0.798913
9	-0.917969	-1.996244	-1.128136e+12	-0.229589
10	-1.927254	-2.152966	-1.426380e+12	-2.003892
11	0.756422	0.751297	-6.406850e+11	0.738550
12	-3.610867	-3.884441	1.531544e+12	-2.944202

In [58]: ▶

1 *##http://zetcode.com/python/prettytable/*  
2  
3 from prettytable import PrettyTable  
4 x = PrettyTable()  
5 x.field\_names = [ "Model", "MSE"]  
6 print('MSE comparisons ')  
7 x.add\_row(["Sklearn SGD",22.2])  
8 x.add\_row(["Sklearn Linear Regression",22.1])  
9 x.add\_row(["Custom SGD (Trial 1)",8.9e+24])  
10 x.add\_row(["Custom SGD (Trial 2)",22.5])  
11 print(x)

MSE comparisons

Model	MSE
Sklearn SGD	22.2
Sklearn Linear Regression	22.1
Custom SGD (Trial 1)	8.9e+24
Custom SGD (Trial 2)	22.5

Conclusion:

- 1. After applying custom SGD to our dataset with learning rate 0.2 and iterations of 500 we see that MSE is closer to the SKlearn SGD MSE 22.2.

Ref :<https://www.kaggle.com/premvardhan/stochasticgradientdescent-implementation-lr-python> (<https://www.kaggle.com/premvardhan/stochasticgradientdescent-implementation-lr-python>),  
<https://www.kaggle.com/tentotheminus9/linear-regression-from-scratch-gradient-descent> (<https://www.kaggle.com/tentotheminus9/linear-regression-from-scratch-gradient-descent>)

In [ ]: ▶

1