

```
In [32]: 1 ## Import Libraries
2 # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
3 from keras.utils import np_utils
4 from keras.datasets import mnist
5 import seaborn as sns
6 import pandas as pd
7 from keras.initializers import RandomNormal
8 from keras import backend as K
9 from keras import optimizers, losses
10 from keras.layers import Dense, Activation
11 from keras.layers import BatchNormalization, Dropout
12 from keras.models import Sequential
```

```
In [33]: 1 ## Load dataset
2 (x_train, y_train), (x_test, y_test) = mnist.load_data()
3
4 ## Loading dataset
5 (x_train, y_train), (x_test, y_test) = mnist.load_data()
6
7 ## network parameters
8 img_rows, img_cols = 28, 28
9 batch_size = 128
10 n_epoch = 15
11 classes = 10
12 ##
13 x_train = x_train.reshape(x_train.shape[0], x_train.shape[1]*x_train.shape[2])
14 x_test = x_test.reshape(x_test.shape[0], x_test.shape[1]*x_test.shape[2])
15 ##
16 print("Number of training examples :", x_train.shape[0], "and each image is of shape (%d)"%(x_train.shape[1]))
17 print("Number of test examples :", x_test.shape[0], "and each image is of shape (%d)"%(x_test.shape[1]))
18
19 ## Normalizing the x_train dataset
20 x_train = x_train.astype('float32')
21 x_test = x_test.astype('float32')
22
23 x_train /= 255
24 x_test /= 255
25
26 print('x_train shape:', x_train.shape)
27 print(x_train.shape[0], 'train samples')
28 print(x_test.shape[0], 'test samples')
29
30 ## one hot encode the target labels
31 y_train = np_utils.to_categorical(y_train, classes)
32 y_test = np_utils.to_categorical(y_test, classes)
33
34 ##
35 input_dim = x_train.shape[1]
36 output_dim = 10
```

```
Number of training examples : 60000 and each image is of shape (784)
Number of test examples : 10000 and each image is of shape (784)
x_train shape: (60000, 784)
60000 train samples
10000 test samples
```

Model1 : 2 Layer

- Without BN + Dropout
 1. number of hidden layers : 2
 2. optimizer : Adam
 3. Activation : Relu
- With BN + Dropout
 1. number of hidden layers : 2
 2. optimizer : Adam
 3. Activation : Relu

Without BN + Dropout

```
In [34]: 1 model = Sequential()
2 model.add(Dense(64, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(stddev=0.050))) ## he initializer sqrt(2/784) = 0.050
3 model.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(stddev=0.176)))## he initializer sqrt(2/64) = 0.0176
4 model.add(Dense(output_dim, activation='softmax'))
5 model.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
=====		
dense_21 (Dense)	(None, 64)	50240
dense_22 (Dense)	(None, 128)	8320
dense_23 (Dense)	(None, 10)	1290
=====		
Total params: 59,850		
Trainable params: 59,850		
Non-trainable params: 0		

```
In [35]: 1 model.compile(optimizer='adam' ,loss='categorical_crossentropy',metrics=['accuracy'] )
```

In [36]:

```

1  ## train the model
2  history = model.fit(x_train,y_train,batch_size=batch_size,epochs=n_epoch,verbose=1,validation_data=[x_test,y_test])

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/15
60000/60000 [=====] - 1s 25us/step - loss: 0.3552 - accuracy: 0.8990 - val_loss: 0.1816 - val_accuracy: 0.9459
Epoch 2/15
60000/60000 [=====] - 1s 19us/step - loss: 0.1466 - accuracy: 0.9573 - val_loss: 0.1400 - val_accuracy: 0.9585
Epoch 3/15
60000/60000 [=====] - 1s 20us/step - loss: 0.1069 - accuracy: 0.9684 - val_loss: 0.0993 - val_accuracy: 0.9684
Epoch 4/15
60000/60000 [=====] - 1s 22us/step - loss: 0.0853 - accuracy: 0.9745 - val_loss: 0.1051 - val_accuracy: 0.9683
Epoch 5/15
60000/60000 [=====] - 2s 27us/step - loss: 0.0685 - accuracy: 0.9794 - val_loss: 0.0847 - val_accuracy: 0.9718
Epoch 6/15
60000/60000 [=====] - 1s 23us/step - loss: 0.0581 - accuracy: 0.9825 - val_loss: 0.0826 - val_accuracy: 0.9750
Epoch 7/15
60000/60000 [=====] - 1s 22us/step - loss: 0.0481 - accuracy: 0.9851 - val_loss: 0.0831 - val_accuracy: 0.9734
Epoch 8/15
60000/60000 [=====] - 1s 20us/step - loss: 0.0401 - accuracy: 0.9876 - val_loss: 0.0736 - val_accuracy: 0.9773
Epoch 9/15
60000/60000 [=====] - 1s 24us/step - loss: 0.0336 - accuracy: 0.9897 - val_loss: 0.0793 - val_accuracy: 0.9760
Epoch 10/15
60000/60000 [=====] - 2s 25us/step - loss: 0.0299 - accuracy: 0.9908 - val_loss: 0.0792 - val_accuracy: 0.9766
Epoch 11/15
60000/60000 [=====] - 1s 24us/step - loss: 0.0250 - accuracy: 0.9924 - val_loss: 0.0767 - val_accuracy: 0.9776
Epoch 12/15
60000/60000 [=====] - 1s 22us/step - loss: 0.0219 - accuracy: 0.9933 - val_loss: 0.0790 - val_accuracy: 0.9772
Epoch 13/15
60000/60000 [=====] - 1s 24us/step - loss: 0.0203 - accuracy: 0.9939 - val_loss: 0.0800 - val_accuracy: 0.9771
Epoch 14/15
60000/60000 [=====] - 1s 22us/step - loss: 0.0163 - accuracy: 0.9948 - val_loss: 0.0875 - val_accuracy: 0.9759
Epoch 15/15
60000/60000 [=====] - 1s 22us/step - loss: 0.0148 - accuracy: 0.9956 - val_loss: 0.0909 - val_accuracy: 0.9746

```

In [37]:

```


1  %matplotlib notebook
2  import matplotlib.pyplot as plt
3  import numpy as np
4  import time
5  # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
6  # https://stackoverflow.com/a/14434334
7  # this function is used to update the plots for each epoch and error
8  def plt_dynamic(x, vy, ty, ax, colors=['b']):
9      ax.plot(x, vy, 'b', label="Validation Loss")
10     ax.plot(x, ty, 'r', label="Train Loss")
11     plt.legend()
12     plt.grid()
13     plt.show()
14     #fig.canvas.draw()

```

```
In [38]: 1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4 fig,ax=plt.subplots(1,1)
5 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
6 # list of epoch numbers
7 x = list(range(1,n_epoch+1))
8 vy = history.history['val_loss']
9 ty = history.history['loss']
10 plt_dynamic(x, vy, ty, ax)
11 result = [[ty[-1],vy[-1],score[1]]]
```

Test score: 0.09085136720940645

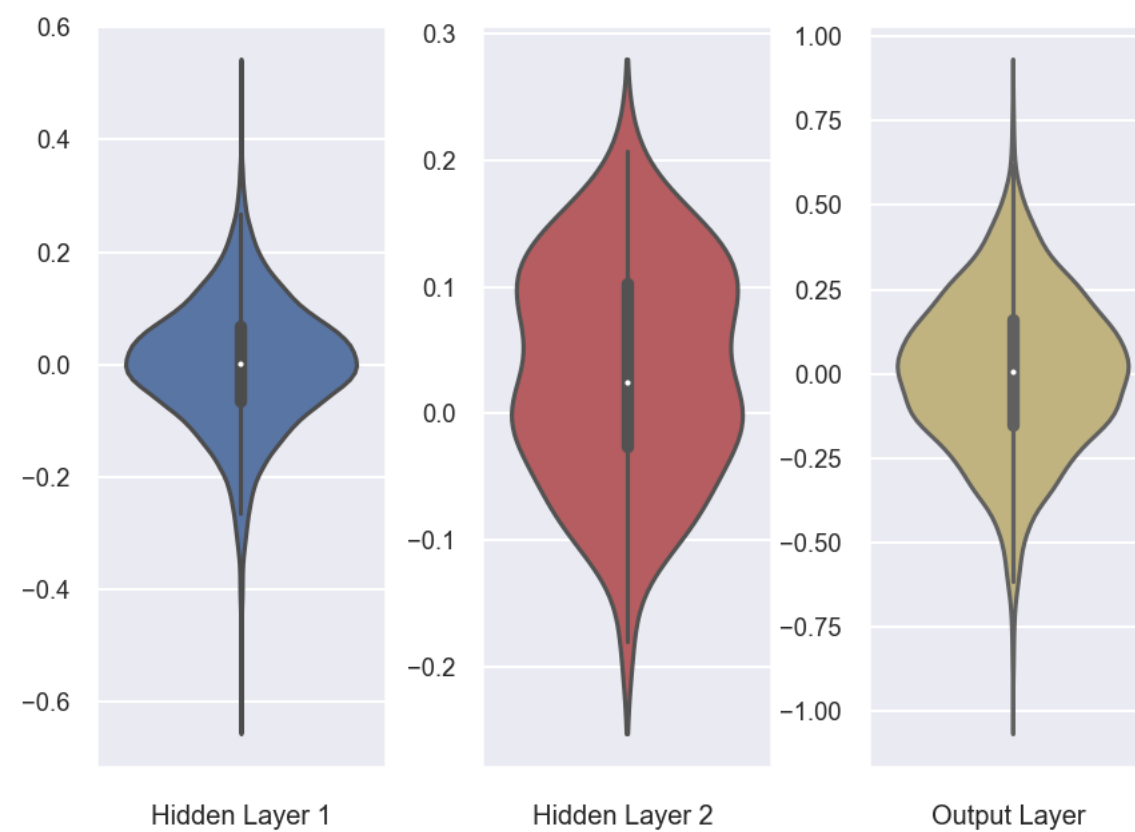
Test accuracy: 0.9746000170707703



```

In [39]: 1 w_after = model.get_weights()
          2
          3 h1_w = w_after[0].flatten().reshape(-1,1)
          4 h2_w = w_after[1].flatten().reshape(-1,1)
          5 out_w = w_after[2].flatten().reshape(-1,1)
          6
          7
          8 fig = plt.figure()
          9 plt.title("Weight matrices after model trained")
         10 plt.subplot(1, 3, 1)
         11 plt.tight_layout(pad=3.0)
         12 ax = sns.violinplot(y=h1_w,color='b')
         13 plt.xlabel('Hidden Layer 1')
         14 plt.subplot(1, 3, 2)
         15 ax = sns.violinplot(y=h2_w, color='r')
         16 plt.xlabel('Hidden Layer 2 ')
         17 plt.subplot(1, 3, 3)
         18 ax = sns.violinplot(y=out_w,color='y')
         19 plt.xlabel('Output Layer ')
         20 plt.show()

```



In [40]: ▶

1


2

```
from sklearn.metrics import confusion_matrix
```

```

In [41]: 1 y_train_predict = model.predict(x_train)
2 y_test_predict = model.predict(x_test)
3
4 c=confusion_matrix(y_train.argmax(axis=1), y_train_predict.argmax(axis=1))
5 normed_c = c.T / c.astype(np.float).sum(axis=1).T
6 df_cm1 = pd.DataFrame(normed_c, range(10), range(10))
7
8 c=confusion_matrix(y_test.argmax(axis=1), y_test_predict.argmax(axis=1))
9 normed_c = c.T / c.astype(np.float).sum(axis=1).T
10 df_cm2 = pd.DataFrame(normed_c, range(10), range(10))
11
12 plt.figure(figsize=(11,4))
13 cmap=sns.light_palette("Gray")
14 labels =[0,1,2,3,4,5,6,7,8,9]
15
16 plt.subplot(1,2,1)
17 sns.set(font_scale=0.8)
18 sns.heatmap(df_cm1,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
19 plt.ylabel('Actual class')
20 plt.xlabel('Predicted class')
21 plt.title('Confusion matrix : Train data')
22
23 plt.subplot(1,2,2)
24 sns.set(font_scale=0.8)
25 sns.heatmap(df_cm2,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
26 plt.ylabel('Actual class')
27 plt.xlabel('Predicted class')
28 plt.title('Confusion matrix : Test data')
29 plt.show()

```



With BN + Dropout

```
In [43]: 1 model = Sequential()
2 model.add(Dense(64, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(stddev=0.050))) ## he initializer sqrt(2/784) = 0.050
3 ## add a batch normalization layer
4 model.add(BatchNormalization())
5 model.add(Dropout(0.25))
6 model.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(stddev=0.176)))## he initializer sqrt(2/64) = 0.0176
7 model.add(Dense(output_dim, activation='softmax'))
8 model.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
=====		
dense_27 (Dense)	(None, 64)	50240
<hr/>		
batch_normalization_5 (Batch Normalization)	(None, 64)	256
<hr/>		
dropout_5 (Dropout)	(None, 64)	0
<hr/>		
dense_28 (Dense)	(None, 128)	8320
<hr/>		
dense_29 (Dense)	(None, 10)	1290
=====		
Total params: 60,106		
Trainable params: 59,978		
Non-trainable params: 128		
<hr/>		

```
In [44]: 1 model.compile(optimizer='adam' ,loss='categorical_crossentropy',metrics=['accuracy'] )
```


In [45]:

```
1  ## train the model
2  history = model.fit(x_train,y_train,batch_size=batch_size,epochs=n_epoch,verbose=1,validation_data=[x_test,y_test])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 2s 34us/step - loss: 0.4044 - accuracy: 0.8772 - val_loss: 0.1741 - val_accuracy: 0.9461

Epoch 2/15

60000/60000 [=====] - 1s 25us/step - loss: 0.1974 - accuracy: 0.9403 - val_loss: 0.1314 - val_accuracy: 0.9622

Epoch 3/15

60000/60000 [=====] - 1s 25us/step - loss: 0.1586 - accuracy: 0.9516 - val_loss: 0.1104 - val_accuracy: 0.9651

Epoch 4/15

60000/60000 [=====] - 2s 26us/step - loss: 0.1365 - accuracy: 0.9574 - val_loss: 0.1057 - val_accuracy: 0.9685

Epoch 5/15

60000/60000 [=====] - 1s 25us/step - loss: 0.1226 - accuracy: 0.9610 - val_loss: 0.0920 - val_accuracy: 0.9718

Epoch 6/15

60000/60000 [=====] - 1s 25us/step - loss: 0.1092 - accuracy: 0.9657 - val_loss: 0.0850 - val_accuracy: 0.9741

Epoch 7/15

60000/60000 [=====] - 2s 30us/step - loss: 0.1052 - accuracy: 0.9672 - val_loss: 0.0849 - val_accuracy: 0.9740

Epoch 8/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0949 - accuracy: 0.9696 - val_loss: 0.0824 - val_accuracy: 0.9749

Epoch 9/15

60000/60000 [=====] - 2s 30us/step - loss: 0.0919 - accuracy: 0.9711 - val_loss: 0.0830 - val_accuracy: 0.9745

Epoch 10/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0868 - accuracy: 0.9720 - val_loss: 0.0784 - val_accuracy: 0.9755

Epoch 11/15

60000/60000 [=====] - 2s 27us/step - loss: 0.0826 - accuracy: 0.9740 - val_loss: 0.0778 - val_accuracy: 0.9765

Epoch 12/15

60000/60000 [=====] - 2s 27us/step - loss: 0.0798 - accuracy: 0.9751 - val_loss: 0.0734 - val_accuracy: 0.9780

Epoch 13/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0777 - accuracy: 0.9744 - val_loss: 0.0758 - val_accuracy: 0.9763

Epoch 14/15

60000/60000 [=====] - 2s 27us/step - loss: 0.0738 - accuracy: 0.9759 - val_loss: 0.0792 - val_accuracy: 0.9750


Epoch 15/15

60000/60000 [=====] - 2s 27us/step - loss: 0.0717 - accuracy: 0.9767 - val_loss: 0.0718 - val_accuracy: 0.9778

```
In [46]: 1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4 fig,ax=plt.subplots(1,1)
5 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
6 # list of epoch numbers
7 x = list(range(1,n_epoch+1))
8 vy = history.history['val_loss']
9 ty = history.history['loss']
10 plt_dynamic(x, vy, ty, ax)
11
12 result.append([ty[-1],vy[-1],score[1]])
```

Test score: 0.07176502713353838

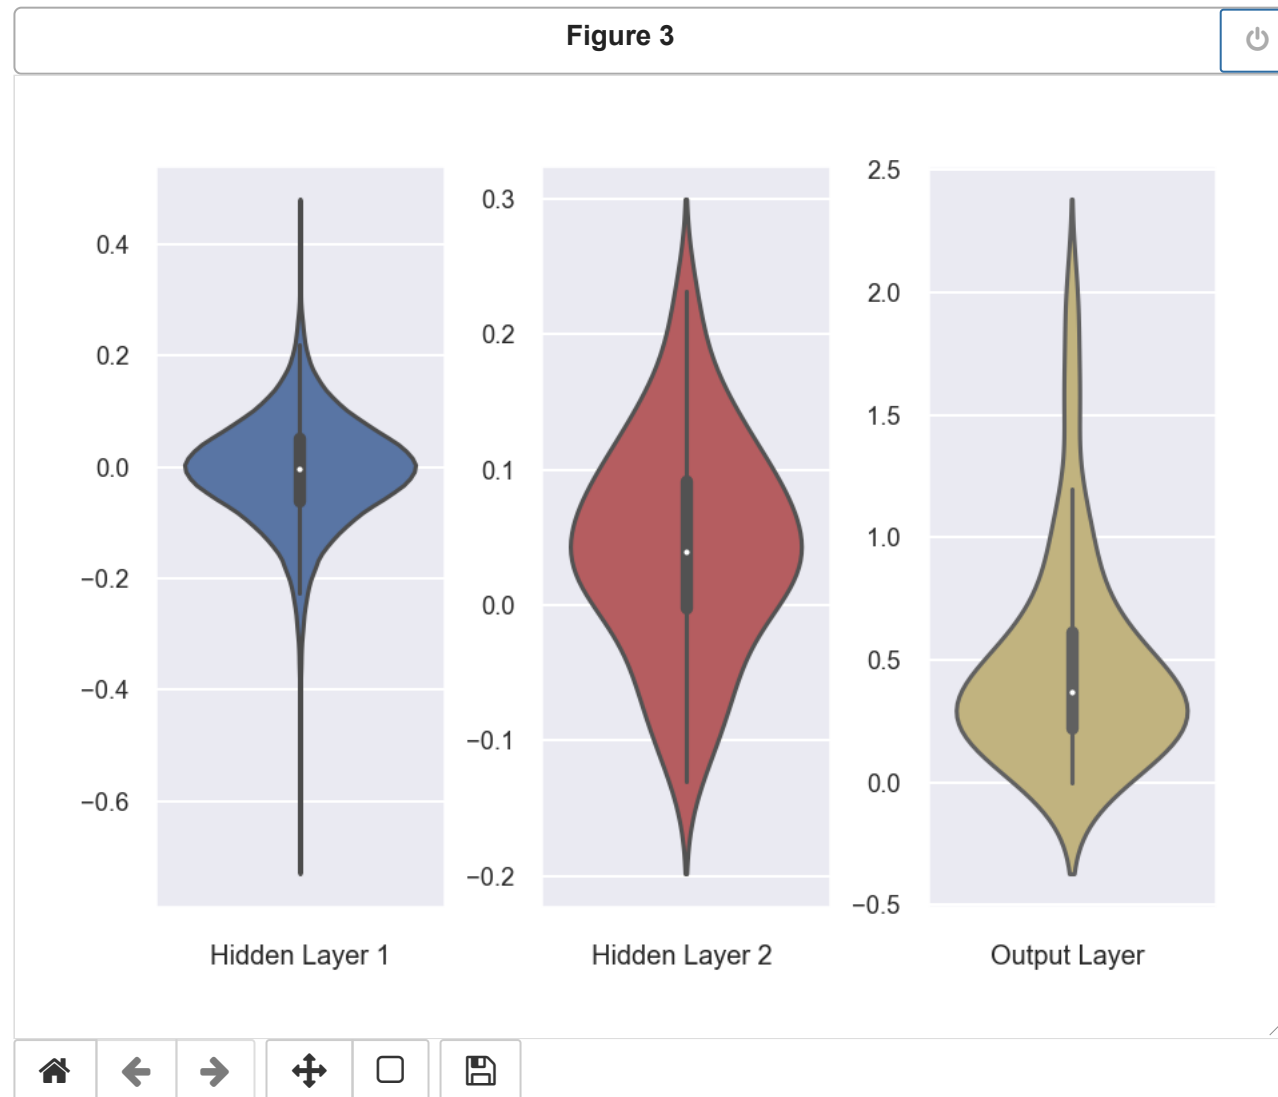
Test accuracy: 0.9778000116348267




```

In [47]: 1 w_after = model.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[3].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.tight_layout(pad=3.0)
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14 plt.subplot(1, 3, 2)
15 ax = sns.violinplot(y=h2_w, color='r')
16 plt.xlabel('Hidden Layer 2 ')
17 plt.subplot(1, 3, 3)
18 ax = sns.violinplot(y=out_w,color='y')
19 plt.xlabel('Output Layer ')
20 plt.show()


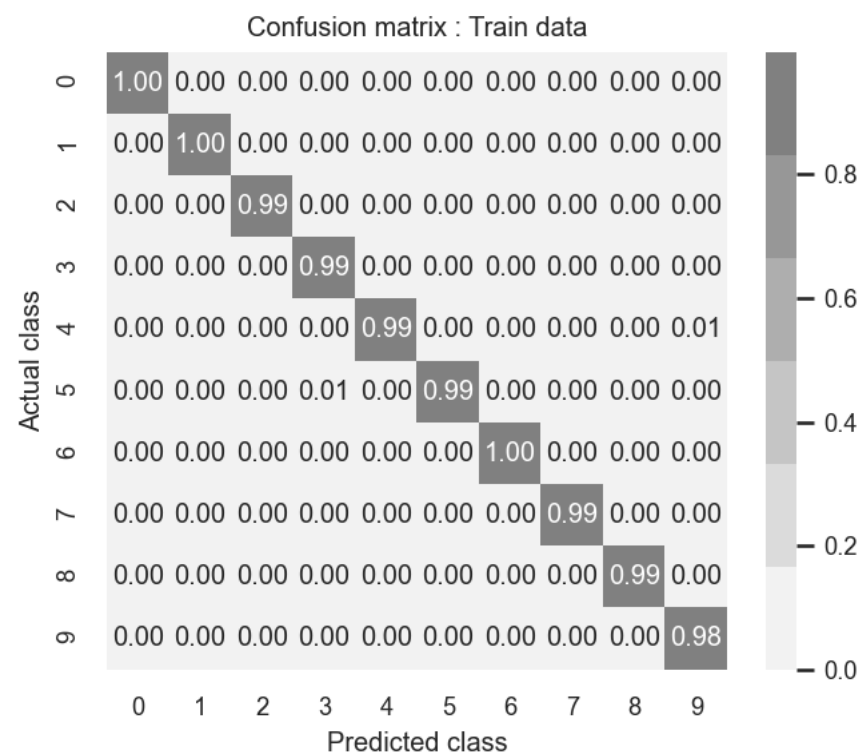
```



```

In [48]: 1 y_train_predict = model.predict(x_train)
2 y_test_predict = model.predict(x_test)
3
4 c=confusion_matrix(y_train.argmax(axis=1), y_train_predict.argmax(axis=1))
5 normed_c = c.T / c.astype(np.float).sum(axis=1).T
6 df_cm1 = pd.DataFrame(normed_c, range(10), range(10))
7
8 c=confusion_matrix(y_test.argmax(axis=1), y_test_predict.argmax(axis=1))
9 normed_c = c.T / c.astype(np.float).sum(axis=1).T
10 df_cm2 = pd.DataFrame(normed_c, range(10), range(10))
11
12 plt.figure(figsize=(11,4))
13 cmap=sns.light_palette("Gray")
14 labels =[0,1,2,3,4,5,6,7,8,9]
15
16 plt.subplot(1,2,1)
17 sns.set(font_scale=0.8)
18 sns.heatmap(df_cm1,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
19 plt.ylabel('Actual class')
20 plt.xlabel('Predicted class')
21 plt.title('Confusion matrix : Train data')
22
23 plt.subplot(1,2,2)
24 sns.set(font_scale=0.8)
25 sns.heatmap(df_cm2,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
26 plt.ylabel('Actual class')
27 plt.xlabel('Predicted class')
28 plt.title('Confusion matrix : Test data')
29 plt.show()

```



Model2 : 3 Layer

- Without BN + Dropout
 1. number of hidden layers : 3
 2. optimizer : Adam
 3. Activation : Relu
- With BN + Dropout
 1. number of hidden layers : 3
 2. optimizer : Adam
 3. Activation : Relu

Without BN + Dropout

```
In [49]: 1 model = Sequential()
2 model.add(Dense(64, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(stddev=0.050))) ## he initializer sqrt(2/784) = 0.050
3 model.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(stddev=0.176)))## he initializer sqrt(2/64) = 0.0176
4 model.add(Dense(256, activation='relu',kernel_initializer=RandomNormal(stddev=0.125)))## he initializer sqrt(2/128) =0.125
5 model.add(Dense(output_dim, activation='softmax'))
6 model.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
=====		
dense_30 (Dense)	(None, 64)	50240
dense_31 (Dense)	(None, 128)	8320
dense_32 (Dense)	(None, 256)	33024
dense_33 (Dense)	(None, 10)	2570
=====		
Total params: 94,154		
Trainable params: 94,154		
Non-trainable params: 0		

```
In [50]: 1 model.compile(optimizer='adam' ,loss='categorical_crossentropy',metrics=['accuracy'] )
```

In [51]:

```
1  ## train the model
2  history = model.fit(x_train,y_train,batch_size=batch_size,epochs=n_epoch,verbose=1,validation_data=[x_test,y_test])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 2s 30us/step - loss: 0.3039 - accuracy: 0.9090 - val_loss: 0.1508 - val_accuracy: 0.9534

Epoch 2/15

60000/60000 [=====] - 2s 25us/step - loss: 0.1217 - accuracy: 0.9631 - val_loss: 0.1101 - val_accuracy: 0.9672

Epoch 3/15

60000/60000 [=====] - 1s 24us/step - loss: 0.0900 - accuracy: 0.9722 - val_loss: 0.1102 - val_accuracy: 0.9670

Epoch 4/15

60000/60000 [=====] - 2s 25us/step - loss: 0.0683 - accuracy: 0.9784 - val_loss: 0.0956 - val_accuracy: 0.9718

Epoch 5/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0544 - accuracy: 0.9826 - val_loss: 0.0962 - val_accuracy: 0.9723

Epoch 6/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0463 - accuracy: 0.9850 - val_loss: 0.0944 - val_accuracy: 0.9734

Epoch 7/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0374 - accuracy: 0.9876 - val_loss: 0.1019 - val_accuracy: 0.9724

Epoch 8/15

60000/60000 [=====] - 2s 31us/step - loss: 0.0335 - accuracy: 0.9884 - val_loss: 0.1000 - val_accuracy: 0.9721

Epoch 9/15

60000/60000 [=====] - 2s 31us/step - loss: 0.0290 - accuracy: 0.9906 - val_loss: 0.0870 - val_accuracy: 0.9765

Epoch 10/15

60000/60000 [=====] - 2s 28us/step - loss: 0.0259 - accuracy: 0.9913 - val_loss: 0.0937 - val_accuracy: 0.9755

Epoch 11/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0222 - accuracy: 0.9924 - val_loss: 0.0913 - val_accuracy: 0.9777

Epoch 12/15

60000/60000 [=====] - 2s 27us/step - loss: 0.0205 - accuracy: 0.9928 - val_loss: 0.0921 - val_accuracy: 0.9766

Epoch 13/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0210 - accuracy: 0.9931 - val_loss: 0.1008 - val_accuracy: 0.9745

Epoch 14/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0177 - accuracy: 0.9937 - val_loss: 0.0964 - val_accuracy: 0.9760


Epoch 15/15

60000/60000 [=====] - 2s 29us/step - loss: 0.0129 - accuracy: 0.9955 - val_loss: 0.1019 - val_accuracy: 0.9758

```
In [52]: 1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4 fig,ax=plt.subplots(1,1)
5 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
6 # list of epoch numbers
7 x = list(range(1,n_epoch+1))
8 vy = history.history['val_loss']
9 ty = history.history['loss']
10 plt_dynamic(x, vy, ty, ax)
11 result.append([ty[-1],vy[-1],score[1]])
```

Test score: 0.10194387945031066


Test accuracy: 0.9757999777793884




```

In [53]: 1 w_after = model.get_weights()
          2
          3 h1_w = w_after[0].flatten().reshape(-1,1)
          4 h2_w = w_after[1].flatten().reshape(-1,1)
          5 h3_w = w_after[2].flatten().reshape(-1,1)
          6 out_w = w_after[3].flatten().reshape(-1,1)
          7
          8
          9 fig = plt.figure()
         10 plt.title("Weight matrices after model trained")
         11 plt.subplot(1, 4, 1)
         12 plt.tight_layout(pad=3.0)
         13 ax = sns.violinplot(y=h1_w,color='b')
         14 plt.xlabel('Hidden Layer 1')
         15 plt.subplot(1, 4, 2)
         16 ax = sns.violinplot(y=h2_w, color='r')
         17 plt.xlabel('Hidden Layer 2 ')
         18 plt.subplot(1, 4, 3)
         19 ax = sns.violinplot(y=h3_w,color='g')
         20 plt.xlabel('Hidden Layer 3 ')
         21 plt.subplot(1, 4, 4)
         22 ax = sns.violinplot(y=out_w,color='y')
         23 plt.xlabel('Output Layer ')
         24 plt.show()



```



```

In [54]: 1 y_train_predict = model.predict(x_train)
2 y_test_predict = model.predict(x_test)
3
4 c=confusion_matrix(y_train.argmax(axis=1), y_train_predict.argmax(axis=1))
5 normed_c = c.T / c.astype(np.float).sum(axis=1).T
6 df_cm1 = pd.DataFrame(normed_c, range(10), range(10))
7
8 c=confusion_matrix(y_test.argmax(axis=1), y_test_predict.argmax(axis=1))
9 normed_c = c.T / c.astype(np.float).sum(axis=1).T
10 df_cm2 = pd.DataFrame(normed_c, range(10), range(10))
11
12 plt.figure(figsize=(11,4))
13 cmap=sns.light_palette("Gray")
14 labels =[0,1,2,3,4,5,6,7,8,9]
15
16 plt.subplot(1,2,1)
17 sns.set(font_scale=0.8)
18 sns.heatmap(df_cm1,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
19 plt.ylabel('Actual class')
20 plt.xlabel('Predicted class')
21 plt.title('Confusion matrix : Train data')
22
23 plt.subplot(1,2,2)
24 sns.set(font_scale=0.8)
25 sns.heatmap(df_cm2,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
26 plt.ylabel('Actual class')
27 plt.xlabel('Predicted class')
28 plt.title('Confusion matrix : Test data')
29 plt.show()

```



With BN + Dropout

```
In [55]: 1 model = Sequential()
2 model.add(Dense(64, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(stddev=0.050))) ## he initializer sqrt(2/784) = 0.050
3 model.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(stddev=0.176)))## he initializer sqrt(2/64) = 0.0176
4 model.add(BatchNormalization())
5 model.add(Dropout(0.25))
6 model.add(Dense(256, activation='relu',kernel_initializer=RandomNormal(stddev=0.125)))## he initializer sqrt(2/128) =0.125
7 model.add(BatchNormalization())
8 model.add(Dropout(0.5))
9 model.add(Dense(output_dim, activation='softmax'))
10 model.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
=====		
dense_34 (Dense)	(None, 64)	50240
dense_35 (Dense)	(None, 128)	8320
batch_normalization_6 (Batch Normalization)	(None, 128)	512
dropout_6 (Dropout)	(None, 128)	0
dense_36 (Dense)	(None, 256)	33024
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dropout_7 (Dropout)	(None, 256)	0
dense_37 (Dense)	(None, 10)	2570
=====		
Total params: 95,690		
Trainable params: 94,922		
Non-trainable params: 768		

```
In [56]: 1 model.compile(optimizer='adam' ,loss='categorical_crossentropy',metrics=['accuracy'] )
```

In [57]:

```

1  ## train the model
2  history = model.fit(x_train,y_train,batch_size=batch_size,epochs=n_epoch,verbose=1,validation_data=[x_test,y_test])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 3s 49us/step - loss: 0.5035 - accuracy: 0.8484 - val_loss: 0.1878 - val_accuracy: 0.9441

Epoch 2/15

60000/60000 [=====] - 2s 41us/step - loss: 0.2099 - accuracy: 0.9378 - val_loss: 0.1373 - val_accuracy: 0.9581

Epoch 3/15

60000/60000 [=====] - 3s 44us/step - loss: 0.1568 - accuracy: 0.9530 - val_loss: 0.1114 - val_accuracy: 0.9670

Epoch 4/15

60000/60000 [=====] - 3s 52us/step - loss: 0.1249 - accuracy: 0.9621 - val_loss: 0.1089 - val_accuracy: 0.9676

Epoch 5/15

60000/60000 [=====] - 3s 47us/step - loss: 0.1061 - accuracy: 0.9679 - val_loss: 0.1072 - val_accuracy: 0.9670

Epoch 6/15

60000/60000 [=====] - 3s 51us/step - loss: 0.0925 - accuracy: 0.9715 - val_loss: 0.0965 - val_accuracy: 0.9694

Epoch 7/15

60000/60000 [=====] - 3s 46us/step - loss: 0.0825 - accuracy: 0.9740 - val_loss: 0.1009 - val_accuracy: 0.9700

Epoch 8/15

60000/60000 [=====] - 3s 48us/step - loss: 0.0741 - accuracy: 0.9766 - val_loss: 0.0964 - val_accuracy: 0.9723

Epoch 9/15

60000/60000 [=====] - 3s 47us/step - loss: 0.0659 - accuracy: 0.9795 - val_loss: 0.0937 - val_accuracy: 0.9737

Epoch 10/15

60000/60000 [=====] - 3s 48us/step - loss: 0.0616 - accuracy: 0.9810 - val_loss: 0.0874 - val_accuracy: 0.9756

Epoch 11/15

60000/60000 [=====] - 3s 47us/step - loss: 0.0551 - accuracy: 0.9825 - val_loss: 0.0957 - val_accuracy: 0.9735

Epoch 12/15

60000/60000 [=====] - 3s 48us/step - loss: 0.0528 - accuracy: 0.9831 - val_loss: 0.0866 - val_accuracy: 0.9768

Epoch 13/15

60000/60000 [=====] - 3s 47us/step - loss: 0.0462 - accuracy: 0.9853 - val_loss: 0.1020 - val_accuracy: 0.9733

Epoch 14/15

60000/60000 [=====] - 3s 47us/step - loss: 0.0449 - accuracy: 0.9853 - val_loss: 0.0939 - val_accuracy: 0.9747

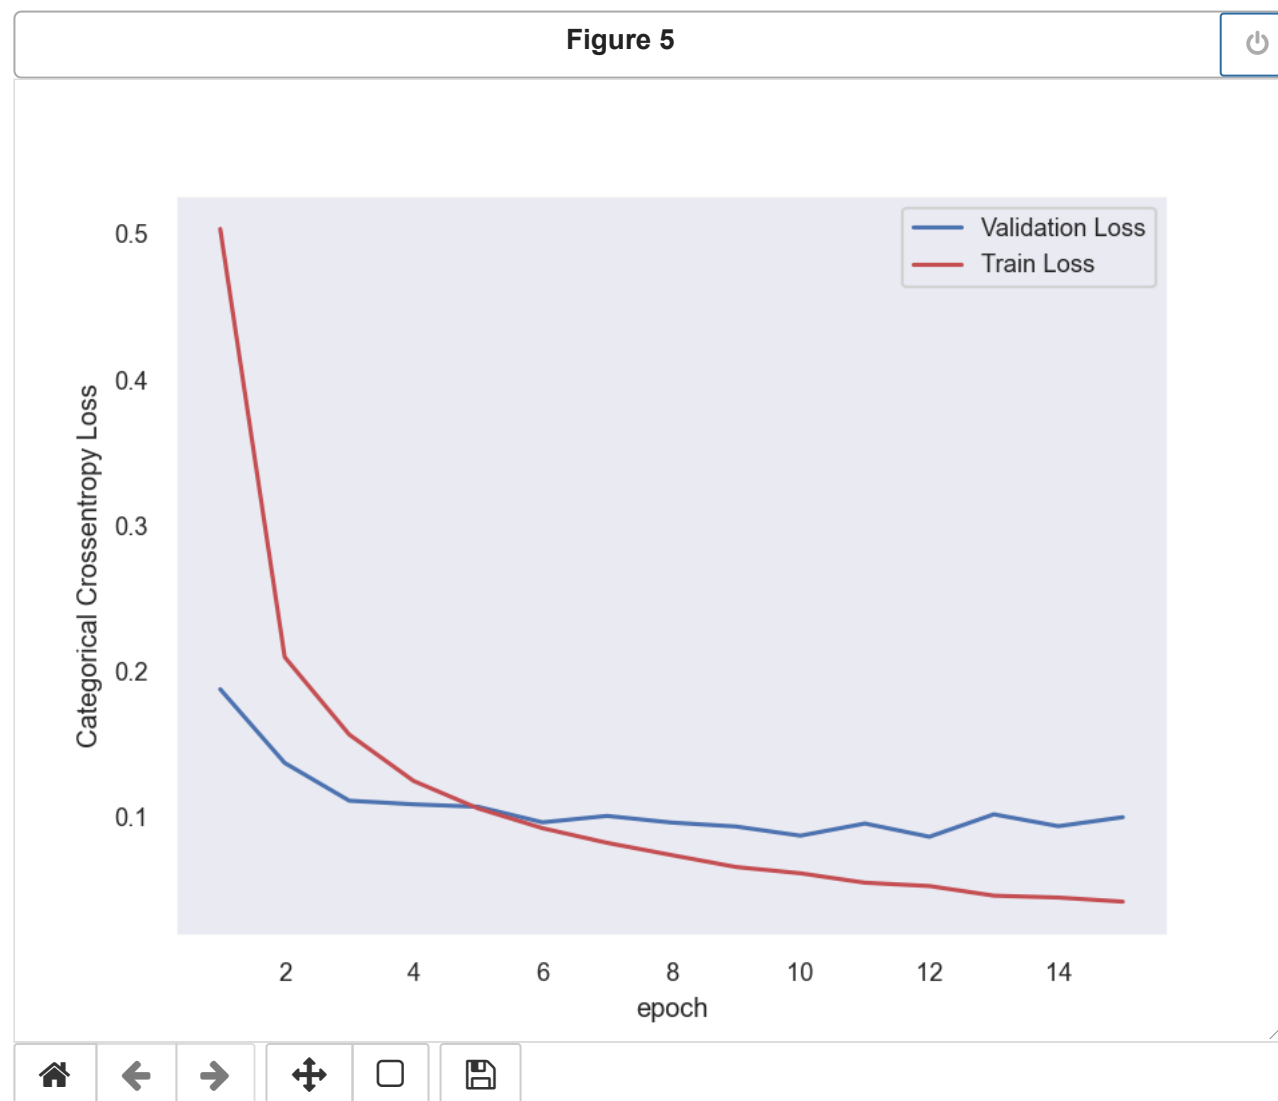
Epoch 15/15

60000/60000 [=====] - 3s 50us/step - loss: 0.0421 - accuracy: 0.9862 - val_loss: 0.1000 - val_accuracy: 0.9739

```
In [58]: 1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4 fig,ax=plt.subplots(1,1)
5 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
6 # list of epoch numbers
7 x = list(range(1,n_epoch+1))
8 vy = history.history['val_loss']
9 ty = history.history['loss']
10 plt_dynamic(x, vy, ty, ax)
11 result.append([ty[-1],vy[-1],score[1]])
```

Test score: 0.10004436637778999


Test accuracy: 0.9739000201225281




```

In [59]: 1 w_after = model.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[1].flatten().reshape(-1,1)
5 h3_w = w_after[4].flatten().reshape(-1,1)
6 out_w = w_after[7].flatten().reshape(-1,1)
7
8
9 fig = plt.figure()
10 plt.title("Weight matrices after model trained")
11 plt.subplot(1, 4, 1)
12 plt.tight_layout(pad=3.0)
13 ax = sns.violinplot(y=h1_w,color='b')
14 plt.xlabel('Hidden Layer 1')
15 plt.subplot(1, 4, 2)
16 ax = sns.violinplot(y=h2_w, color='r')
17 plt.xlabel('Hidden Layer 2 ')
18 plt.subplot(1, 4, 3)
19 ax = sns.violinplot(y=h3_w,color='g')
20 plt.xlabel('Hidden Layer 3 ')
21 plt.subplot(1, 4, 4)
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()



```




```

In [60]: 1 y_train_predict = model.predict(x_train)
2 y_test_predict = model.predict(x_test)
3
4 c=confusion_matrix(y_train.argmax(axis=1), y_train_predict.argmax(axis=1))
5 normed_c = c.T / c.astype(np.float).sum(axis=1).T
6 df_cm1 = pd.DataFrame(normed_c, range(10), range(10))
7
8 c=confusion_matrix(y_test.argmax(axis=1), y_test_predict.argmax(axis=1))
9 normed_c = c.T / c.astype(np.float).sum(axis=1).T
10 df_cm2 = pd.DataFrame(normed_c, range(10), range(10))
11
12 plt.figure(figsize=(11,4))
13 cmap=sns.light_palette("Gray")
14 labels =[0,1,2,3,4,5,6,7,8,9]
15
16 plt.subplot(1,2,1)
17 sns.set(font_scale=0.8)
18 sns.heatmap(df_cm1,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
19 plt.ylabel('Actual class')
20 plt.xlabel('Predicted class')
21 plt.title('Confusion matrix : Train data')
22
23 plt.subplot(1,2,2)
24 sns.set(font_scale=0.8)
25 sns.heatmap(df_cm2,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
26 plt.ylabel('Actual class')
27 plt.xlabel('Predicted class')
28 plt.title('Confusion matrix : Test data')
29 plt.show()

```



Model3 : 5 Layer

- Without BN + Dropout
 1. number of hidden layers : 5
 2. optimizer : Adam
 3. Activation : Relu
- With BN + Dropout
 1. number of hidden layers : 5
 2. optimizer : Adam
 3. Activation : Relu

Without BN + Dropout

```
In [64]: 1 model = Sequential()
2 model.add(Dense(32, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(stddev=0.125)))## he initializer sqrt(2/784) =0.050
3 model.add(Dense(64, activation='relu',kernel_initializer=RandomNormal(stddev=0.25))) ## he initializer sqrt(2/32) = 0.25
4 model.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(stddev=0.176)))## he initializer sqrt(2/64) = 0.17
5 model.add(Dense(256, activation='relu',kernel_initializer=RandomNormal(stddev=0.125)))## he initializer sqrt(2/128) =0.125
6 model.add(Dense(512, activation='relu',kernel_initializer=RandomNormal(stddev=0.088)))## he initializer sqrt(2/256) =0.088
7 model.add(Dense(output_dim, activation='softmax'))
8 model.summary()
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
=====		
dense_44 (Dense)	(None, 32)	25120
dense_45 (Dense)	(None, 64)	2112
dense_46 (Dense)	(None, 128)	8320
dense_47 (Dense)	(None, 256)	33024
dense_48 (Dense)	(None, 512)	131584
dense_49 (Dense)	(None, 10)	5130
=====		
Total params: 205,290		
Trainable params: 205,290		
Non-trainable params: 0		

```
In [65]: 1 model.compile(optimizer='adam' ,loss='categorical_crossentropy',metrics=['accuracy'] )
```

In [66]:

```
1  ## train the model
2  history = model.fit(x_train,y_train,batch_size=batch_size,epochs=n_epoch,verbose=1,validation_data=[x_test,y_test])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 3s 52us/step - loss: 0.3952 - accuracy: 0.8754 - val_loss: 0.2033 - val_accuracy: 0.9386

Epoch 2/15

60000/60000 [=====] - 3s 45us/step - loss: 0.1660 - accuracy: 0.9489 - val_loss: 0.1436 - val_accuracy: 0.9562

Epoch 3/15

60000/60000 [=====] - 3s 54us/step - loss: 0.1186 - accuracy: 0.9630 - val_loss: 0.1268 - val_accuracy: 0.9611

Epoch 4/15

60000/60000 [=====] - 3s 48us/step - loss: 0.0954 - accuracy: 0.9689 - val_loss: 0.1218 - val_accuracy: 0.9618

Epoch 5/15

60000/60000 [=====] - 3s 52us/step - loss: 0.0781 - accuracy: 0.9744 - val_loss: 0.1234 - val_accuracy: 0.9631

Epoch 6/15

60000/60000 [=====] - 3s 55us/step - loss: 0.0639 - accuracy: 0.9795 - val_loss: 0.1328 - val_accuracy: 0.9635

Epoch 7/15

60000/60000 [=====] - 3s 51us/step - loss: 0.0568 - accuracy: 0.9815 - val_loss: 0.1223 - val_accuracy: 0.9665

Epoch 8/15

60000/60000 [=====] - 3s 47us/step - loss: 0.0476 - accuracy: 0.9839 - val_loss: 0.1333 - val_accuracy: 0.9661

Epoch 9/15

60000/60000 [=====] - 4s 59us/step - loss: 0.0400 - accuracy: 0.9867 - val_loss: 0.1369 - val_accuracy: 0.9676

Epoch 10/15

60000/60000 [=====] - 4s 67us/step - loss: 0.0389 - accuracy: 0.9871 - val_loss: 0.1249 - val_accuracy: 0.9691

Epoch 11/15

60000/60000 [=====] - 3s 51us/step - loss: 0.0373 - accuracy: 0.9878 - val_loss: 0.1369 - val_accuracy: 0.9671

Epoch 12/15

60000/60000 [=====] - 3s 49us/step - loss: 0.0285 - accuracy: 0.9910 - val_loss: 0.1389 - val_accuracy: 0.9679

Epoch 13/15

60000/60000 [=====] - 3s 48us/step - loss: 0.0280 - accuracy: 0.9908 - val_loss: 0.1415 - val_accuracy: 0.9696

Epoch 14/15

60000/60000 [=====] - 3s 51us/step - loss: 0.0272 - accuracy: 0.9909 - val_loss: 0.1554 - val_accuracy: 0.9671

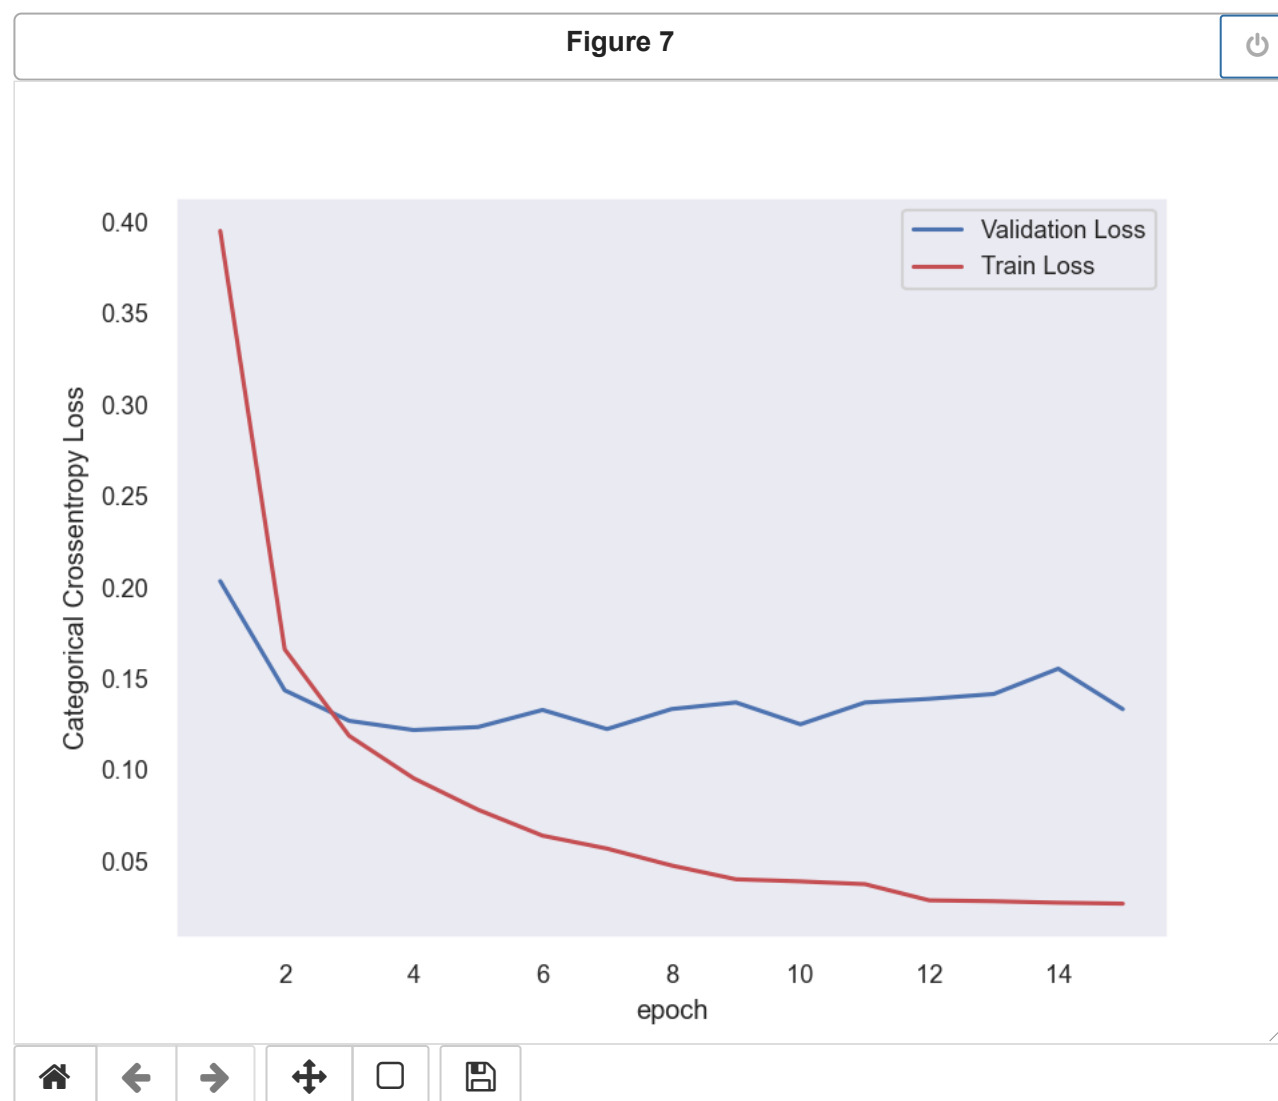
Epoch 15/15

60000/60000 [=====] - 3s 48us/step - loss: 0.0267 - accuracy: 0.9910 - val_loss: 0.1332 - val_accuracy: 0.9698

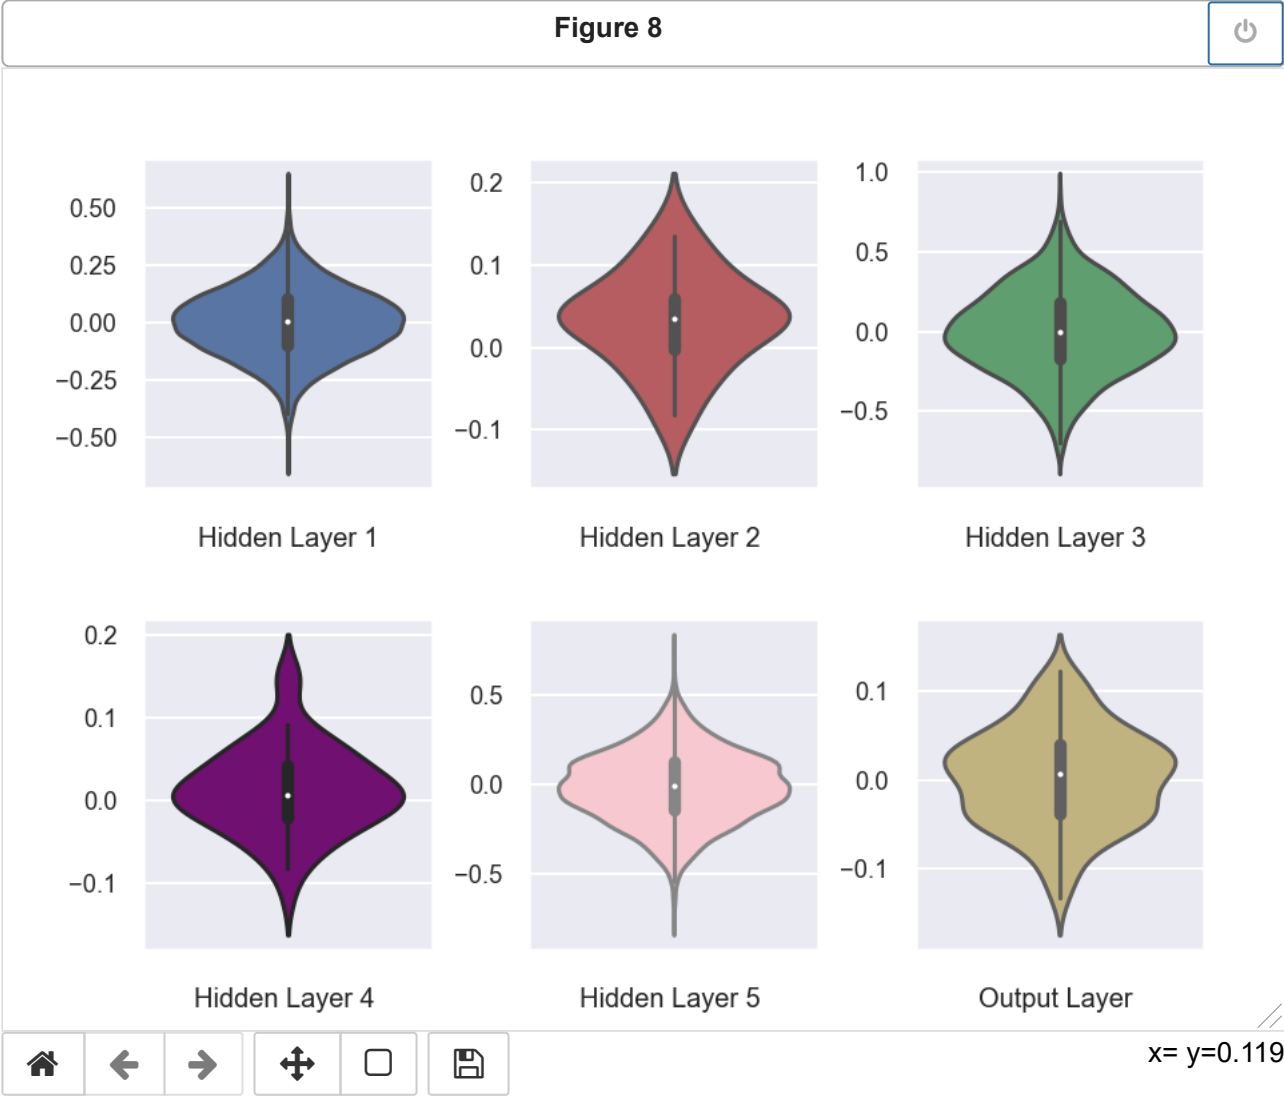
```
In [67]: 1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4 fig,ax=plt.subplots(1,1)
5 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
6 # list of epoch numbers
7 x = list(range(1,n_epoch+1))
8 vy = history.history['val_loss']
9 ty = history.history['loss']
10 plt_dynamic(x, vy, ty, ax)
11 result.append([ty[-1],vy[-1],score[1]])
```

Test score: 0.13321278729955158

Test accuracy: 0.9697999954223633




```
In [68]: 1 w_after = model.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[1].flatten().reshape(-1,1)
5 h3_w = w_after[2].flatten().reshape(-1,1)
6 h4_w = w_after[3].flatten().reshape(-1,1)
7 h5_w = w_after[4].flatten().reshape(-1,1)
8 out_w = w_after[5].flatten().reshape(-1,1)
9
10
11 fig = plt.figure()
12 plt.title("Weight matrices after model trained")
13 plt.subplot(2, 3, 1)
14 plt.tight_layout(pad=3.0)
15 ax = sns.violinplot(y=h1_w,color='b')
16 plt.xlabel('Hidden Layer 1')
17 plt.subplot(2, 3, 2)
18 ax = sns.violinplot(y=h2_w, color='r')
19 plt.xlabel('Hidden Layer 2 ')
20 plt.subplot(2, 3,3)
21 ax = sns.violinplot(y=h3_w,color='g')
22 plt.xlabel('Hidden Layer 3 ')
23 plt.subplot(2, 3, 4)
24 ax = sns.violinplot(y=h4_w,color='purple')
25 plt.xlabel('Hidden Layer 4 ')
26 plt.subplot(2, 3, 5)
27 ax = sns.violinplot(y=h5_w,color='pink')
28 plt.xlabel('Hidden Layer 5 ')
29 plt.subplot(2, 3, 6)
30 ax = sns.violinplot(y=out_w,color='y')
31 plt.xlabel('Output Layer ')
32 plt.show()
```




```

In [69]: 1 y_train_predict = model.predict(x_train)
2 y_test_predict = model.predict(x_test)
3
4 c=confusion_matrix(y_train.argmax(axis=1), y_train_predict.argmax(axis=1))
5 normed_c = c.T / c.astype(np.float).sum(axis=1).T
6 df_cm1 = pd.DataFrame(normed_c, range(10), range(10))
7
8 c=confusion_matrix(y_test.argmax(axis=1), y_test_predict.argmax(axis=1))
9 normed_c = c.T / c.astype(np.float).sum(axis=1).T
10 df_cm2 = pd.DataFrame(normed_c, range(10), range(10))
11
12 plt.figure(figsize=(11,4))
13 cmap=sns.light_palette("Gray")
14 labels =[0,1,2,3,4,5,6,7,8,9]
15
16 plt.subplot(1,2,1)
17 sns.set(font_scale=0.8)
18 sns.heatmap(df_cm1,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
19 plt.ylabel('Actual class')
20 plt.xlabel('Predicted class')
21 plt.title('Confusion matrix : Train data')
22
23 plt.subplot(1,2,2)
24 sns.set(font_scale=0.8)
25 sns.heatmap(df_cm2,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
26 plt.ylabel('Actual class')
27 plt.xlabel('Predicted class')
28 plt.title('Confusion matrix : Test data')
29 plt.show()

```

Figure 9



With BN + Dropout

```
In [70]: 1 model = Sequential()
2 model.add(Dense(32, activation='relu',input_shape=(input_dim,),kernel_initializer=RandomNormal(stddev=0.125)))## he initializer sqrt(2/784) =0.050
3 model.add(Dense(64, activation='relu',kernel_initializer=RandomNormal(stddev=0.25))) ## he initializer sqrt(2/32) = 0.25
4 model.add(BatchNormalization())
5 model.add(Dropout(0.25))
6 model.add(Dense(128, activation='relu',kernel_initializer=RandomNormal(stddev=0.176)))## he initializer sqrt(2/64) = 0.17
7 model.add(Dense(256, activation='relu',kernel_initializer=RandomNormal(stddev=0.125)))## he initializer sqrt(2/128) =0.125
8 model.add(BatchNormalization())
9 model.add(Dropout(0.5))
10 model.add(Dense(512, activation='relu',kernel_initializer=RandomNormal(stddev=0.088)))## he initializer sqrt(2/256) =0.088
11 model.add(Dense(output_dim, activation='softmax'))
12 model.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
=====		
dense_50 (Dense)	(None, 32)	25120
dense_51 (Dense)	(None, 64)	2112
batch_normalization_8 (Batch Normalization)	(None, 64)	256
dropout_8 (Dropout)	(None, 64)	0
dense_52 (Dense)	(None, 128)	8320
dense_53 (Dense)	(None, 256)	33024
batch_normalization_9 (Batch Normalization)	(None, 256)	1024
dropout_9 (Dropout)	(None, 256)	0
dense_54 (Dense)	(None, 512)	131584
dense_55 (Dense)	(None, 10)	5130
=====		
Total params: 206,570		
Trainable params: 205,930		
Non-trainable params: 640		

```
In [71]: 1 model.compile(optimizer='adam' ,loss='categorical_crossentropy',metrics=['accuracy'] )
```

In [72]:

```
1  ## train the model
2  history = model.fit(x_train,y_train,batch_size=batch_size,epochs=n_epoch,verbose=1,validation_data=[x_test,y_test])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 4s 73us/step - loss: 0.7466 - accuracy: 0.7624 - val_loss: 0.2717 - val_accuracy: 0.9167

Epoch 2/15

60000/60000 [=====] - 5s 80us/step - loss: 0.3306 - accuracy: 0.9003 - val_loss: 0.1889 - val_accuracy: 0.9417

Epoch 3/15

60000/60000 [=====] - 4s 65us/step - loss: 0.2539 - accuracy: 0.9233 - val_loss: 0.1583 - val_accuracy: 0.9512

Epoch 4/15

60000/60000 [=====] - 4s 70us/step - loss: 0.2121 - accuracy: 0.9358 - val_loss: 0.1483 - val_accuracy: 0.9551

Epoch 5/15

60000/60000 [=====] - 4s 65us/step - loss: 0.1865 - accuracy: 0.9438 - val_loss: 0.1416 - val_accuracy: 0.9575

Epoch 6/15

60000/60000 [=====] - 5s 85us/step - loss: 0.1720 - accuracy: 0.9475 - val_loss: 0.1279 - val_accuracy: 0.9596

Epoch 7/15

60000/60000 [=====] - 4s 72us/step - loss: 0.1587 - accuracy: 0.9523 - val_loss: 0.1204 - val_accuracy: 0.9645

Epoch 8/15

60000/60000 [=====] - 4s 75us/step - loss: 0.1439 - accuracy: 0.9573 - val_loss: 0.1230 - val_accuracy: 0.9649

Epoch 9/15

60000/60000 [=====] - 5s 76us/step - loss: 0.1333 - accuracy: 0.9595 - val_loss: 0.1088 - val_accuracy: 0.9692

Epoch 10/15

60000/60000 [=====] - 4s 74us/step - loss: 0.1306 - accuracy: 0.9607 - val_loss: 0.1081 - val_accuracy: 0.9680

Epoch 11/15

60000/60000 [=====] - 4s 72us/step - loss: 0.1232 - accuracy: 0.9626 - val_loss: 0.1096 - val_accuracy: 0.9689

Epoch 12/15

60000/60000 [=====] - 4s 74us/step - loss: 0.1118 - accuracy: 0.9665 - val_loss: 0.1232 - val_accuracy: 0.9657

Epoch 13/15

60000/60000 [=====] - 5s 81us/step - loss: 0.1083 - accuracy: 0.9669 - val_loss: 0.1106 - val_accuracy: 0.9689

Epoch 14/15

60000/60000 [=====] - 4s 69us/step - loss: 0.1038 - accuracy: 0.9689 - val_loss: 0.1089 - val_accuracy: 0.9688

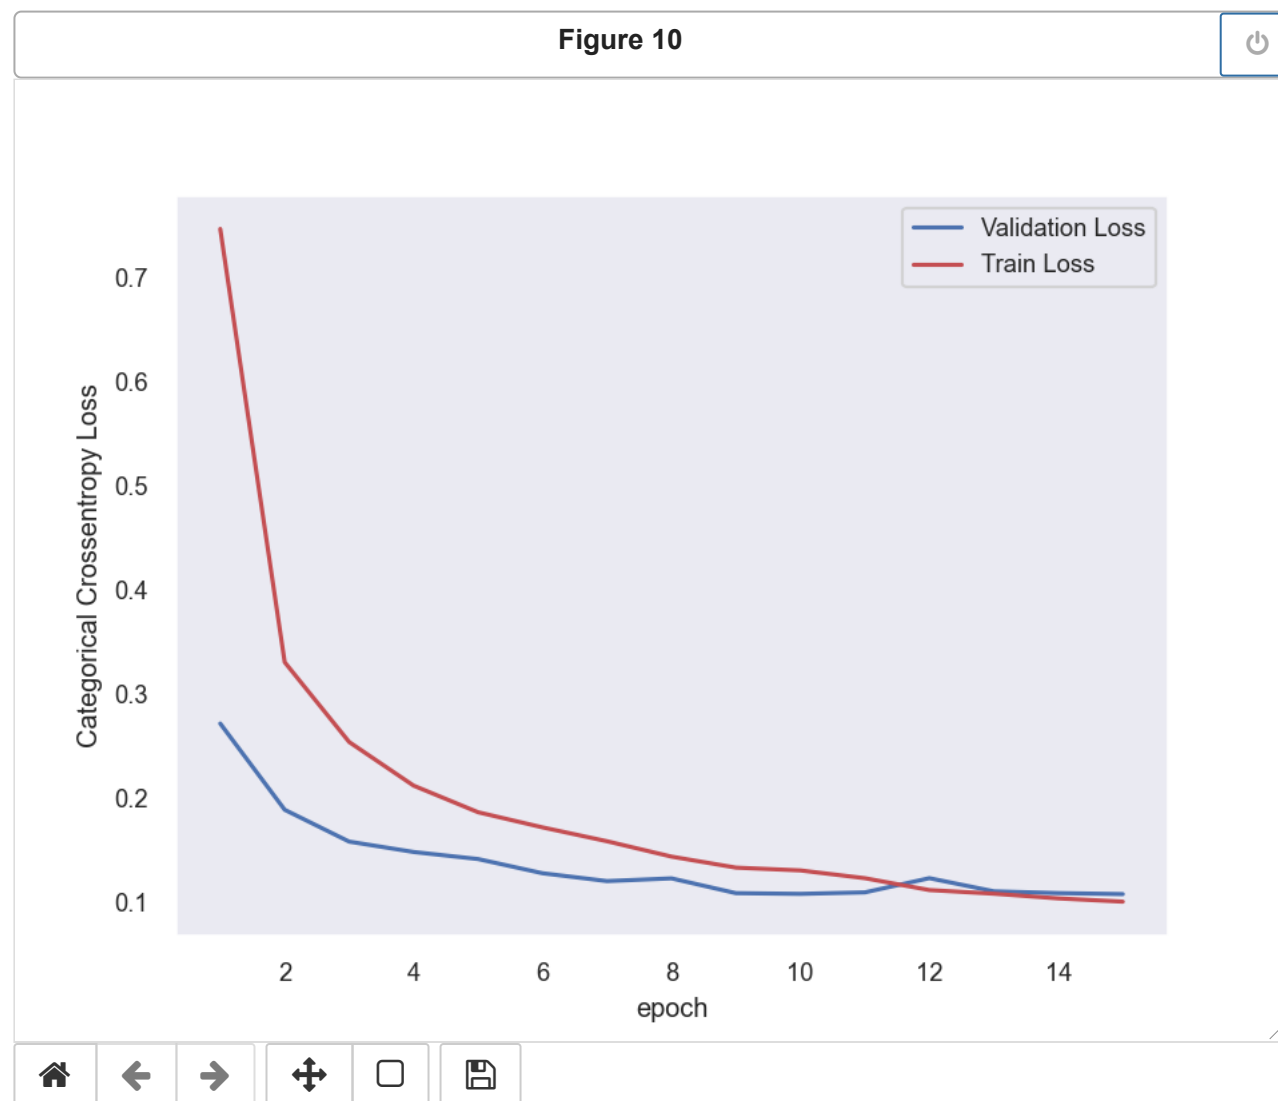
Epoch 15/15

60000/60000 [=====] - 5s 77us/step - loss: 0.1007 - accuracy: 0.9690 - val_loss: 0.1079 - val_accuracy: 0.9703

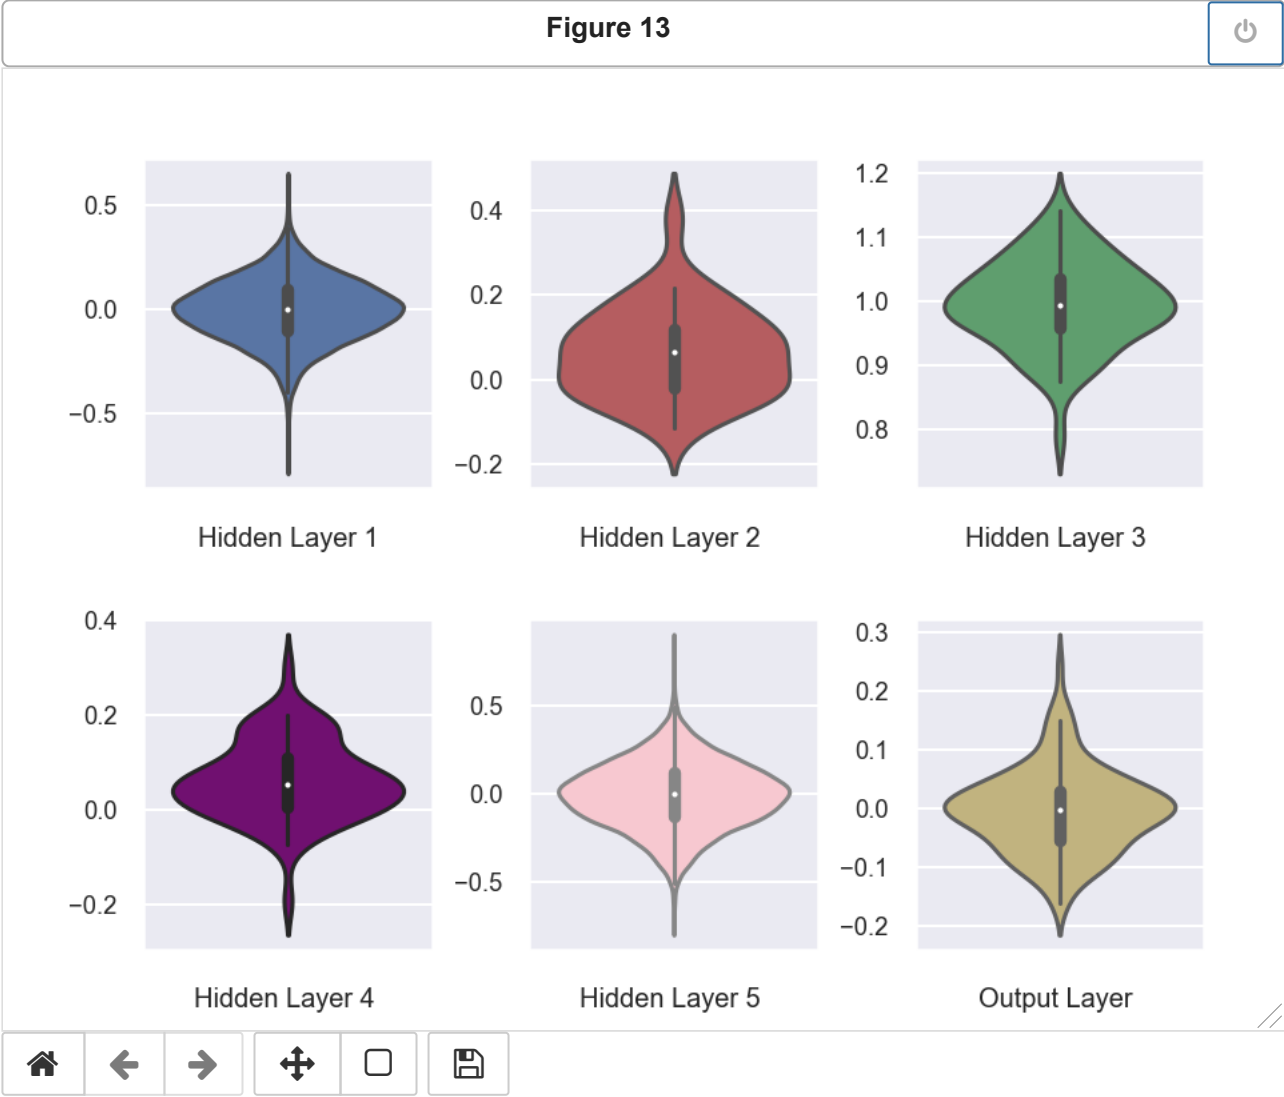
```
In [73]: 1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4 fig,ax=plt.subplots(1,1)
5 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
6 # list of epoch numbers
7 x = list(range(1,n_epoch+1))
8 vy = history.history['val_loss']
9 ty = history.history['loss']
10 plt_dynamic(x, vy, ty, ax)
11 result.append([ty[-1],vy[-1],score[1]])
```

Test score: 0.10790942553719506

Test accuracy: 0.970300018787384






```
In [77]: 1 w_after = model.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[1].flatten().reshape(-1,1)
5 h3_w = w_after[4].flatten().reshape(-1,1)
6 h4_w = w_after[5].flatten().reshape(-1,1)
7 h5_w = w_after[8].flatten().reshape(-1,1)
8 out_w = w_after[9].flatten().reshape(-1,1)
9
10
11 fig = plt.figure()
12 plt.title("Weight matrices after model trained")
13 plt.subplot(2, 3, 1)
14 plt.tight_layout(pad=3.0)
15 ax = sns.violinplot(y=h1_w,color='b')
16 plt.xlabel('Hidden Layer 1')
17 plt.subplot(2, 3, 2)
18 ax = sns.violinplot(y=h2_w, color='r')
19 plt.xlabel('Hidden Layer 2 ')
20 plt.subplot(2, 3,3)
21 ax = sns.violinplot(y=h3_w,color='g')
22 plt.xlabel('Hidden Layer 3 ')
23 plt.subplot(2, 3, 4)
24 ax = sns.violinplot(y=h4_w,color='purple')
25 plt.xlabel('Hidden Layer 4 ')
26 plt.subplot(2, 3, 5)
27 ax = sns.violinplot(y=h5_w,color='pink')
28 plt.xlabel('Hidden Layer 5 ')
29 plt.subplot(2, 3, 6)
30 ax = sns.violinplot(y=out_w,color='y')
31 plt.xlabel('Output Layer ')
32 plt.show()
```



```

In [75]: 1 y_train_predict = model.predict(x_train)
2 y_test_predict = model.predict(x_test)
3
4 c=confusion_matrix(y_train.argmax(axis=1), y_train_predict.argmax(axis=1))
5 normed_c = c.T / c.astype(np.float).sum(axis=1).T
6 df_cm1 = pd.DataFrame(normed_c, range(10), range(10))
7
8 c=confusion_matrix(y_test.argmax(axis=1), y_test_predict.argmax(axis=1))
9 normed_c = c.T / c.astype(np.float).sum(axis=1).T
10 df_cm2 = pd.DataFrame(normed_c, range(10), range(10))
11
12 plt.figure(figsize=(11,4))
13 cmap=sns.light_palette("Gray")
14 labels =[0,1,2,3,4,5,6,7,8,9]
15
16 plt.subplot(1,2,1)
17 sns.set(font_scale=0.8)
18 sns.heatmap(df_cm1,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
19 plt.ylabel('Actual class')
20 plt.xlabel('Predicted class')
21 plt.title('Confusion matrix : Train data')
22
23 plt.subplot(1,2,2)
24 sns.set(font_scale=0.8)
25 sns.heatmap(df_cm2,annot=True,fmt=".2f",xticklabels=labels,yticklabels=labels,cmap=cmap)
26 plt.ylabel('Actual class')
27 plt.xlabel('Predicted class')
28 plt.title('Confusion matrix : Test data')
29 plt.show()

```



```

In [78]: 1  ##http://zetcode.com/python/prettytable/
2
3  from prettytable import PrettyTable
4
5  x = PrettyTable()
6  x.field_names = ["Model", "Number of Layer", 'Batch Normalization and Drop out', 'Train-loss', 'Test-Loss', "Test-Accuracy"]
7
8  x.add_row(["Model 1", 2, 'no', round(result[0][0], 2), round(result[0][1], 2), round(result[0][2], 2)])
9  x.add_row(["Model 1", 2, 'yes', round(result[1][0], 2), round(result[1][1], 2), round(result[1][2], 2)])
10 x.add_row(["Model 2", 3, 'no', round(result[2][0], 2), round(result[2][1], 2), round(result[2][2], 2)])
11 x.add_row(["Model 2", 3, 'yes', round(result[3][0], 2), round(result[3][1], 2), round(result[3][2], 2)])
12 x.add_row(["Model 3", 5, 'no', round(result[4][0], 2), round(result[4][1], 2), round(result[4][2], 2)])
13 x.add_row(["Model 3", 5, 'yes', round(result[5][0], 2), round(result[5][1], 2), round(result[5][2], 2)])
14 print(x)

```

Model	Number of Layer	Batch Normalization and Drop out	Train-loss	Test-Loss	Test-Accuracy
Model 1	2	no	0.01	0.09	0.97
Model 1	2	yes	0.07	0.07	0.98
Model 2	3	no	0.01	0.1	0.98
Model 2	3	yes	0.04	0.1	0.97
Model 3	5	no	0.03	0.13	0.97
Model 3	5	yes	0.1	0.11	0.97

OBSERVATIONS :

- * Since Mnist is a simple dataset we get optimum accuracy with 2 layers itself.
- * As the number of layers increases model converges faster at early epochs.
- * Adding drop out and batch normalization layers avoids overfitting on data.
- * Layer weights are not too small or too large to cause slow convergence .
- * Number of true positives on test data improved from layer 2 to layer 5

```

In [ ]: 1

```