

```
In [1]: 1 # import keras
2 # from keras.datasets import cifar10
3 # from keras.models import Model, Sequential
4 # from keras.layers import Dense, Dropout, Flatten, Input, AveragePooling2D, merge, Activation
5 # from keras.layers import Conv2D, MaxPooling2D, BatchNormalization
6 # from keras.layers import Concatenate
7 # from keras.optimizers import Adam
8 from tensorflow.keras import models, layers
9 from tensorflow.keras.models import Model
10 from tensorflow.keras.layers import BatchNormalization, Activation, Flatten
11 from tensorflow.keras.optimizers import Adam
```

```
In [2]: 1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [3]: 1 from keras import backend as K
2 import tensorflow as tf
```

```
In [4]: 1 # ! pip install tensorflow
```

```
In [5]: 1 import keras
2 config = tf.compat.v1.Session(config=tf.compat.v1.ConfigProto(log_device_placement=True, device_count = {'GPU': 1 , 'CPU': 56}))
3 keras.backend.set_session(config)
```

Device mapping:

/job:localhost/replica:0/task:0/device:GPU:0 -> device: 0, name: Tesla T4, pci bus id: 0000:00:04.0, compute capability: 7.5

```
In [6]: 1 # this part will prevent tensorflow to allocate all the available GPU Memory
2 # backend
3 import tensorflow as tf
```

```
In [7]: 1 # Hyperparameters
2 batch_size = 128
3 num_classes = 10
4 epochs = 10
5 l = 40
6 num_filter = 12
7 compression = 0.5
8 dropout_rate = 0.2
```

```
In [8]: 1 # Load CIFAR10 Data
2 (X_train, y_train), (X_test, y_test) = tf.keras.datasets.cifar10.load_data()
3 img_height, img_width, channel = X_train.shape[1],X_train.shape[2],X_train.shape[3]
4
5 # convert to one hot encoding
6 y_train = tf.keras.utils.to_categorical(y_train, num_classes)
7 y_test = tf.keras.utils.to_categorical(y_test, num_classes)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> (<https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>)  
 170500096/170498071 [=====] - 4s 0us/step

```
In [9]: 1 X_train.shape
```

Out[9]: (50000, 32, 32, 3)

```
In [10]: 1 X_test.shape
```

Out[10]: (10000, 32, 32, 3)

```
In [11]: 1 # Dense Block
2 def denseblock(input, num_filter = 12, dropout_rate = 0.2):
3     global compression
4     temp = input
5     for _ in range(1):
6         BatchNorm = layers.BatchNormalization()(temp)
7         relu = layers.Activation('relu')(BatchNorm)
8         Conv2D_3_3 = layers.Conv2D(int(num_filter*compression), (3,3), use_bias=False ,padding='same')(relu)
9         if dropout_rate>0:
10             Conv2D_3_3 = layers.Dropout(dropout_rate)(Conv2D_3_3)
11         concat = layers.Concatenate(axis=-1)([temp,Conv2D_3_3])
12
13         temp = concat
14
15     return temp
16
17 ## transition Block
18 def transition(input, num_filter = 12, dropout_rate = 0.2):
19     global compression
20     BatchNorm = layers.BatchNormalization()(input)
21     relu = layers.Activation('relu')(BatchNorm)
22     Conv2D_BottleNeck = layers.Conv2D(int(num_filter*compression), (1,1), use_bias=False ,padding='same')(relu)
23     if dropout_rate>0:
24         Conv2D_BottleNeck = layers.Dropout(dropout_rate)(Conv2D_BottleNeck)
25     avg = layers.AveragePooling2D(pool_size=(2,2))(Conv2D_BottleNeck)
26     return avg
27
28 #output Layer
29 def output_layer(input):
30     global compression
31     BatchNorm = layers.BatchNormalization()(input)
32     relu = layers.Activation('relu')(BatchNorm)
33     AvgPooling = layers.AveragePooling2D(pool_size=(2,2))(relu)
34     flat = layers.Flatten()(AvgPooling)
35     output = layers.Dense(num_classes, activation='softmax')(flat)
36     return output
```

```
In [12]: ▶ 1 num_filter = 12
2 dropout_rate = 0.2
3 l = 12
4 input = layers.Input(shape=(img_height, img_width, channel,))
5 First_Conv2D = layers.Conv2D(num_filter, (3,3), use_bias=False ,padding='same')(input)
6
7 First_Block = denseblock(First_Conv2D, num_filter, dropout_rate)
8 First_Transition = transition(First_Block, num_filter, dropout_rate)
9
10 Second_Block = denseblock(First_Transition, num_filter, dropout_rate)
11 Second_Transition = transition(Second_Block, num_filter, dropout_rate)
12
13 Third_Block = denseblock(Second_Transition, num_filter, dropout_rate)
14 Third_Transition = transition(Third_Block, num_filter, dropout_rate)
15
16 Last_Block = denseblock(Third_Transition, num_filter, dropout_rate)
17 output = output_layer(Last_Block)
```

```
In [ ]: ▶ 1 #https://arxiv.org/pdf/1608.06993.pdf
2 from IPython.display import IFrame, YouTubeVideo
3 YouTubeVideo(id='-W6y8xnd--U', width=600)
```

Out[9]:

```
Model: "model"
```

[illegible]

262

```
1 # determine Loss function and Optimizer
2 model.compile(loss='categorical_crossentropy',
3               optimizer=Adam(),
4               metrics=['accuracy'])
```

```
In [ ]: 1 model.fit(X_train, y_train,
2             batch_size=batch_size,
3             epochs=epochs,
4             verbose=1,
5             validation_data=(X_test, y_test))
```

```
Epoch 1/10
391/391 [=====] - 78s 102ms/step - loss: 1.7187 - accuracy: 0.3524 - val_loss: 1.6062 - val_accuracy: 0.4232
Epoch 2/10
391/391 [=====] - 37s 94ms/step - loss: 1.3829 - accuracy: 0.4883 - val_loss: 1.3495 - val_accuracy: 0.5169
Epoch 3/10
391/391 [=====] - 38s 96ms/step - loss: 1.2379 - accuracy: 0.5463 - val_loss: 1.1956 - val_accuracy: 0.5768
Epoch 4/10
391/391 [=====] - 38s 97ms/step - loss: 1.1237 - accuracy: 0.5931 - val_loss: 1.2230 - val_accuracy: 0.5891
Epoch 5/10
391/391 [=====] - 38s 97ms/step - loss: 1.0488 - accuracy: 0.6229 - val_loss: 1.5108 - val_accuracy: 0.5358
Epoch 6/10
391/391 [=====] - 38s 97ms/step - loss: 0.9915 - accuracy: 0.6443 - val_loss: 1.0291 - val_accuracy: 0.6479
Epoch 7/10
391/391 [=====] - 38s 97ms/step - loss: 0.9515 - accuracy: 0.6586 - val_loss: 1.1797 - val_accuracy: 0.6007
Epoch 8/10
391/391 [=====] - 38s 97ms/step - loss: 0.9201 - accuracy: 0.6681 - val_loss: 1.0407 - val_accuracy: 0.6418
Epoch 9/10
391/391 [=====] - 38s 97ms/step - loss: 0.8911 - accuracy: 0.6808 - val_loss: 0.9578 - val_accuracy: 0.6730
Epoch 10/10
391/391 [=====] - 38s 97ms/step - loss: 0.8642 - accuracy: 0.6904 - val_loss: 0.9221 - val_accuracy: 0.6791
```

Out[13]: <tensorflow.python.keras.callbacks.History at 0x7f879036de10>

```
In [ ]: 1 # Test the model
2 score = model.evaluate(X_test, y_test, verbose=1)
3 print('Test loss:', score[0])
4 print('Test accuracy:', score[1])
```

```
313/313 [=====] - 3s 8ms/step - loss: 0.9221 - accuracy: 0.6791
Test loss: 0.9220625758171082
Test accuracy: 0.679099977016449
```

```
In [ ]: 1 # Save the trained weights in to .h5 format
2 model.save_weights("DNST_model.h5")
3 print("Saved model to disk")
```

Saved model to disk

```
In [ ]: 1 del model
```

## CNN on CIFR Assignment:

1. Please visit this link to access the state-of-art DenseNet code for reference - DenseNet - cifar10 notebook link
2. You need to create a copy of this and "retrain" this model to achieve 90+ test accuracy.
3. You cannot use DropOut layers.
4. You MUST use Image Augmentation Techniques.
5. You cannot use an already trained model as a beginning points, you have to initilize as your own

6. You cannot run the program for more than 300 Epochs, and it should be clear from your log, that you have only used 300 Epochs
7. You cannot use test images for training the model.
8. You cannot change the general architecture of DenseNet (which means you must use Dense Block, Transition and Output blocks as mentioned in the code)
9. You are free to change Convolution types (e.g. from 3x3 normal convolution to Depthwise Separable, etc)
10. You cannot have more than 1 Million parameters in total
11. You are free to move the code from Keras to Tensorflow, Pytorch, MXNET etc.
12. You can use any optimization algorithm you need.
13. You can checkpoint your model and retrain the model from that checkpoint so that no need of training the model from first if you lost at any epoch while training. You can directly load that model and Train from that epoch.

```
In [14]: 1 import matplotlib.pyplot as plt
        2 import numpy as np
```

```
In [15]: 1 # Load CIFAR10 Data
        2 (X_train, y_train), (X_test, y_test) = tf.keras.datasets.cifar10.load_data()
        3 img_height, img_width, channel = X_train.shape[1],X_train.shape[2],X_train.shape[3]
```

```
In [16]: 1 ## visualizing the classes
        2 fig = plt.figure(figsize=(20,11))
        3 classes = ['airplane', 'automobile', 'peacock', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
        4 for i in range(10):
        5     plt.subplot(1,10,i+1)
        6     plt.imshow(X_train[np.where(y_train==i)[0][0]])
        7     plt.title(classes[i])
```



```
In [17]: 1 ##### scale the pixels
        2 ## https://machinelearningmastery.com/how-to-develop-a-cnn-from-scratch-for-cifar-10-photo-classification/
        3 # scale pixels
        4 def scale_pixels(train, test):
        5     # convert from integers to floats
        6     train_norm = train.astype('float32')
        7     test_norm = test.astype('float32')
        8     # normalize to range 0-1
        9     train_norm = train_norm / 255
       10     test_norm = test_norm / 255
       11     # return normalized images
       12     return train_norm, test_norm
```

```
In [18]: 1 X_train,X_test = scale_pixels(X_train,X_test)
```

```
In [19]: 1 # convert to one hot encoding
2 y_train = tf.keras.utils.to_categorical(y_train, num_classes)
3 y_test = tf.keras.utils.to_categorical(y_test, num_classes)
```

let's change the dropout\_rate to 0, alter strides to (5,5) and number of filters

```
In [20]: 1 # Dense Block
2 def denseblock(input, num_filter = 12, dropout_rate = 0):
3     global compression
4     temp = input
5     for _ in range(1):
6         BatchNorm = layers.BatchNormalization()(temp)
7         relu = layers.Activation('relu')(BatchNorm)
8         Conv2D_5_5 = layers.Conv2D(int(num_filter*compression), (5,5), kernel_initializer='he_uniform', use_bias=False, padding='same')(relu)
9         if dropout_rate>0:
10             Conv2D_5_5 = layers.Dropout(dropout_rate)(Conv2D_5_5)
11         concat = layers.Concatenate(axis=-1)([temp, Conv2D_5_5])
12
13         temp = concat
14
15     return temp
16
17 ## transition Block
18 def transition(input, num_filter = 12, dropout_rate = 0):
19     global compression
20     BatchNorm = layers.BatchNormalization()(input)
21     relu = layers.Activation('relu')(BatchNorm)
22     Conv2D_BottleNeck = layers.Conv2D(int(num_filter*compression), (5,5), kernel_initializer='he_uniform', use_bias=False, padding='same')(relu)
23     if dropout_rate>0:
24         Conv2D_BottleNeck = layers.Dropout(dropout_rate)(Conv2D_BottleNeck)
25     avg = layers.AveragePooling2D(pool_size=(2,2))(Conv2D_BottleNeck)
26     return avg
27
28 #output Layer
29 def output_layer(input):
30     global compression
31     BatchNorm = layers.BatchNormalization()(input)
32     relu = layers.Activation('relu')(BatchNorm)
33     AvgPooling = layers.AveragePooling2D(pool_size=(2,2))(relu)
34     flat = layers.Flatten()(AvgPooling)
35     output = layers.Dense(num_classes, activation='softmax')(flat)
36     return output
```

```
In [21]: ▶ 1 num_filter = 10
2 dropout_rate = 0
3 l = 12
4 input = layers.Input(shape=(img_height, img_width, channel))
5 First_Conv2D = layers.Conv2D(num_filter, (5,5), use_bias=False ,padding='same')(input)
6 Batchnorm = layers.BatchNormalization()(First_Conv2D)
7
8 First_Block = denseblock(Batchnorm,32, dropout_rate)
9 First_Transition = transition(First_Block, num_filter, dropout_rate)
10
11 Second_Block = denseblock(First_Transition, 16, dropout_rate)
12 Second_Transition = transition(Second_Block, num_filter, dropout_rate)
13
14 Third_Block = denseblock(Second_Transition, num_filter, dropout_rate)
15 Third_Transition = transition(Third_Block, num_filter, dropout_rate)
16
17 Last_Block = denseblock(Third_Transition, num_filter, dropout_rate)
18 output = output_layer(Last_Block)
```



```
In [22]: 1 model = Model(inputs=[input],outputs=[output])
        2 model.summary()
```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 32, 32, 3)]	0	
conv2d_52 (Conv2D)	(None, 32, 32, 10)	750	input_2[0][0]
batch_normalization_52 (Batch Normalization)	(None, 32, 32, 10)	40	conv2d_52[0][0]
batch_normalization_53 (Batch Normalization)	(None, 32, 32, 10)	40	batch_normalization_52[0][0]
activation_52 (Activation)	(None, 32, 32, 10)	0	batch_normalization_53[0][0]
conv2d_53 (Conv2D)	(None, 32, 32, 16)	4000	activation_52[0][0]
concatenate_48 (Concatenate)	(None, 32, 32, 26)	0	batch_normalization_52[0][0] conv2d_53[0][0]
batch_normalization_54 (Batch Normalization)	(None, 32, 32, 26)	104	concatenate_48[0][0]
activation_53 (Activation)	(None, 32, 32, 26)	0	batch_normalization_54[0][0]
conv2d_54 (Conv2D)	(None, 32, 32, 16)	10400	activation_53[0][0]
concatenate_49 (Concatenate)	(None, 32, 32, 42)	0	concatenate_48[0][0] conv2d_54[0][0]
batch_normalization_55 (Batch Normalization)	(None, 32, 32, 42)	168	concatenate_49[0][0]
activation_54 (Activation)	(None, 32, 32, 42)	0	batch_normalization_55[0][0]
conv2d_55 (Conv2D)	(None, 32, 32, 16)	16800	activation_54[0][0]
concatenate_50 (Concatenate)	(None, 32, 32, 58)	0	concatenate_49[0][0] conv2d_55[0][0]
batch_normalization_56 (Batch Normalization)	(None, 32, 32, 58)	232	concatenate_50[0][0]
activation_55 (Activation)	(None, 32, 32, 58)	0	batch_normalization_56[0][0]
conv2d_56 (Conv2D)	(None, 32, 32, 16)	23200	activation_55[0][0]
concatenate_51 (Concatenate)	(None, 32, 32, 74)	0	concatenate_50[0][0] conv2d_56[0][0]
batch_normalization_57 (Batch Normalization)	(None, 32, 32, 74)	296	concatenate_51[0][0]
activation_56 (Activation)	(None, 32, 32, 74)	0	batch_normalization_57[0][0]
conv2d_57 (Conv2D)	(None, 32, 32, 16)	29600	activation_56[0][0]
concatenate_52 (Concatenate)	(None, 32, 32, 90)	0	concatenate_51[0][0] conv2d_57[0][0]

batch_normalization_58 (BatchNo	(None, 32, 32, 90)	360	concatenate_52[0][0]
activation_57 (Activation)	(None, 32, 32, 90)	0	batch_normalization_58[0][0]
conv2d_58 (Conv2D)	(None, 32, 32, 16)	36000	activation_57[0][0]
concatenate_53 (Concatenate)	(None, 32, 32, 106)	0	concatenate_52[0][0] conv2d_58[0][0]
batch_normalization_59 (BatchNo	(None, 32, 32, 106)	424	concatenate_53[0][0]
activation_58 (Activation)	(None, 32, 32, 106)	0	batch_normalization_59[0][0]
conv2d_59 (Conv2D)	(None, 32, 32, 16)	42400	activation_58[0][0]
concatenate_54 (Concatenate)	(None, 32, 32, 122)	0	concatenate_53[0][0] conv2d_59[0][0]
batch_normalization_60 (BatchNo	(None, 32, 32, 122)	488	concatenate_54[0][0]
activation_59 (Activation)	(None, 32, 32, 122)	0	batch_normalization_60[0][0]
conv2d_60 (Conv2D)	(None, 32, 32, 16)	48800	activation_59[0][0]
concatenate_55 (Concatenate)	(None, 32, 32, 138)	0	concatenate_54[0][0] conv2d_60[0][0]
batch_normalization_61 (BatchNo	(None, 32, 32, 138)	552	concatenate_55[0][0]
activation_60 (Activation)	(None, 32, 32, 138)	0	batch_normalization_61[0][0]
conv2d_61 (Conv2D)	(None, 32, 32, 16)	55200	activation_60[0][0]
concatenate_56 (Concatenate)	(None, 32, 32, 154)	0	concatenate_55[0][0] conv2d_61[0][0]
batch_normalization_62 (BatchNo	(None, 32, 32, 154)	616	concatenate_56[0][0]
activation_61 (Activation)	(None, 32, 32, 154)	0	batch_normalization_62[0][0]
conv2d_62 (Conv2D)	(None, 32, 32, 16)	61600	activation_61[0][0]
concatenate_57 (Concatenate)	(None, 32, 32, 170)	0	concatenate_56[0][0] conv2d_62[0][0]
batch_normalization_63 (BatchNo	(None, 32, 32, 170)	680	concatenate_57[0][0]
activation_62 (Activation)	(None, 32, 32, 170)	0	batch_normalization_63[0][0]
conv2d_63 (Conv2D)	(None, 32, 32, 16)	68000	activation_62[0][0]
concatenate_58 (Concatenate)	(None, 32, 32, 186)	0	concatenate_57[0][0] conv2d_63[0][0]
batch_normalization_64 (BatchNo	(None, 32, 32, 186)	744	concatenate_58[0][0]
activation_63 (Activation)	(None, 32, 32, 186)	0	batch_normalization_64[0][0]
conv2d_64 (Conv2D)	(None, 32, 32, 16)	74400	activation_63[0][0]

concatenate_59 (Concatenate)	(None, 32, 32, 202)	0	concatenate_58[0][0] conv2d_64[0][0]
batch_normalization_65 (BatchNo	(None, 32, 32, 202)	808	concatenate_59[0][0]
activation_64 (Activation)	(None, 32, 32, 202)	0	batch_normalization_65[0][0]
conv2d_65 (Conv2D)	(None, 32, 32, 5)	25250	activation_64[0][0]
average_pooling2d_4 (AveragePoo	(None, 16, 16, 5)	0	conv2d_65[0][0]
batch_normalization_66 (BatchNo	(None, 16, 16, 5)	20	average_pooling2d_4[0][0]
activation_65 (Activation)	(None, 16, 16, 5)	0	batch_normalization_66[0][0]
conv2d_66 (Conv2D)	(None, 16, 16, 8)	1000	activation_65[0][0]
concatenate_60 (Concatenate)	(None, 16, 16, 13)	0	average_pooling2d_4[0][0] conv2d_66[0][0]
batch_normalization_67 (BatchNo	(None, 16, 16, 13)	52	concatenate_60[0][0]
activation_66 (Activation)	(None, 16, 16, 13)	0	batch_normalization_67[0][0]
conv2d_67 (Conv2D)	(None, 16, 16, 8)	2600	activation_66[0][0]
concatenate_61 (Concatenate)	(None, 16, 16, 21)	0	concatenate_60[0][0] conv2d_67[0][0]
batch_normalization_68 (BatchNo	(None, 16, 16, 21)	84	concatenate_61[0][0]
activation_67 (Activation)	(None, 16, 16, 21)	0	batch_normalization_68[0][0]
conv2d_68 (Conv2D)	(None, 16, 16, 8)	4200	activation_67[0][0]
concatenate_62 (Concatenate)	(None, 16, 16, 29)	0	concatenate_61[0][0] conv2d_68[0][0]
batch_normalization_69 (BatchNo	(None, 16, 16, 29)	116	concatenate_62[0][0]
activation_68 (Activation)	(None, 16, 16, 29)	0	batch_normalization_69[0][0]
conv2d_69 (Conv2D)	(None, 16, 16, 8)	5800	activation_68[0][0]
concatenate_63 (Concatenate)	(None, 16, 16, 37)	0	concatenate_62[0][0] conv2d_69[0][0]
batch_normalization_70 (BatchNo	(None, 16, 16, 37)	148	concatenate_63[0][0]
activation_69 (Activation)	(None, 16, 16, 37)	0	batch_normalization_70[0][0]
conv2d_70 (Conv2D)	(None, 16, 16, 8)	7400	activation_69[0][0]
concatenate_64 (Concatenate)	(None, 16, 16, 45)	0	concatenate_63[0][0] conv2d_70[0][0]
batch_normalization_71 (BatchNo	(None, 16, 16, 45)	180	concatenate_64[0][0]

activation_70 (Activation)	(None, 16, 16, 45)	0	batch_normalization_71[0][0]
conv2d_71 (Conv2D)	(None, 16, 16, 8)	9000	activation_70[0][0]
concatenate_65 (Concatenate)	(None, 16, 16, 53)	0	concatenate_64[0][0] conv2d_71[0][0]
batch_normalization_72 (BatchNo	(None, 16, 16, 53)	212	concatenate_65[0][0]
activation_71 (Activation)	(None, 16, 16, 53)	0	batch_normalization_72[0][0]
conv2d_72 (Conv2D)	(None, 16, 16, 8)	10600	activation_71[0][0]
concatenate_66 (Concatenate)	(None, 16, 16, 61)	0	concatenate_65[0][0] conv2d_72[0][0]
batch_normalization_73 (BatchNo	(None, 16, 16, 61)	244	concatenate_66[0][0]
activation_72 (Activation)	(None, 16, 16, 61)	0	batch_normalization_73[0][0]
conv2d_73 (Conv2D)	(None, 16, 16, 8)	12200	activation_72[0][0]
concatenate_67 (Concatenate)	(None, 16, 16, 69)	0	concatenate_66[0][0] conv2d_73[0][0]
batch_normalization_74 (BatchNo	(None, 16, 16, 69)	276	concatenate_67[0][0]
activation_73 (Activation)	(None, 16, 16, 69)	0	batch_normalization_74[0][0]
conv2d_74 (Conv2D)	(None, 16, 16, 8)	13800	activation_73[0][0]
concatenate_68 (Concatenate)	(None, 16, 16, 77)	0	concatenate_67[0][0] conv2d_74[0][0]
batch_normalization_75 (BatchNo	(None, 16, 16, 77)	308	concatenate_68[0][0]
activation_74 (Activation)	(None, 16, 16, 77)	0	batch_normalization_75[0][0]
conv2d_75 (Conv2D)	(None, 16, 16, 8)	15400	activation_74[0][0]
concatenate_69 (Concatenate)	(None, 16, 16, 85)	0	concatenate_68[0][0] conv2d_75[0][0]
batch_normalization_76 (BatchNo	(None, 16, 16, 85)	340	concatenate_69[0][0]
activation_75 (Activation)	(None, 16, 16, 85)	0	batch_normalization_76[0][0]
conv2d_76 (Conv2D)	(None, 16, 16, 8)	17000	activation_75[0][0]
concatenate_70 (Concatenate)	(None, 16, 16, 93)	0	concatenate_69[0][0] conv2d_76[0][0]
batch_normalization_77 (BatchNo	(None, 16, 16, 93)	372	concatenate_70[0][0]
activation_76 (Activation)	(None, 16, 16, 93)	0	batch_normalization_77[0][0]
conv2d_77 (Conv2D)	(None, 16, 16, 8)	18600	activation_76[0][0]
concatenate_71 (Concatenate)	(None, 16, 16, 101)	0	concatenate_70[0][0]

conv2d\_77[0][0]

batch_normalization_78 (BatchNo	(None, 16, 16, 101)	404	concatenate_71[0][0]
activation_77 (Activation)	(None, 16, 16, 101)	0	batch_normalization_78[0][0]
conv2d_78 (Conv2D)	(None, 16, 16, 5)	12625	activation_77[0][0]
average_pooling2d_5 (AveragePoo	(None, 8, 8, 5)	0	conv2d_78[0][0]
batch_normalization_79 (BatchNo	(None, 8, 8, 5)	20	average_pooling2d_5[0][0]
activation_78 (Activation)	(None, 8, 8, 5)	0	batch_normalization_79[0][0]
conv2d_79 (Conv2D)	(None, 8, 8, 5)	625	activation_78[0][0]
concatenate_72 (Concatenate)	(None, 8, 8, 10)	0	average_pooling2d_5[0][0] conv2d_79[0][0]
batch_normalization_80 (BatchNo	(None, 8, 8, 10)	40	concatenate_72[0][0]
activation_79 (Activation)	(None, 8, 8, 10)	0	batch_normalization_80[0][0]
conv2d_80 (Conv2D)	(None, 8, 8, 5)	1250	activation_79[0][0]
concatenate_73 (Concatenate)	(None, 8, 8, 15)	0	concatenate_72[0][0] conv2d_80[0][0]
batch_normalization_81 (BatchNo	(None, 8, 8, 15)	60	concatenate_73[0][0]
activation_80 (Activation)	(None, 8, 8, 15)	0	batch_normalization_81[0][0]
conv2d_81 (Conv2D)	(None, 8, 8, 5)	1875	activation_80[0][0]
concatenate_74 (Concatenate)	(None, 8, 8, 20)	0	concatenate_73[0][0] conv2d_81[0][0]
batch_normalization_82 (BatchNo	(None, 8, 8, 20)	80	concatenate_74[0][0]
activation_81 (Activation)	(None, 8, 8, 20)	0	batch_normalization_82[0][0]
conv2d_82 (Conv2D)	(None, 8, 8, 5)	2500	activation_81[0][0]
concatenate_75 (Concatenate)	(None, 8, 8, 25)	0	concatenate_74[0][0] conv2d_82[0][0]
batch_normalization_83 (BatchNo	(None, 8, 8, 25)	100	concatenate_75[0][0]
activation_82 (Activation)	(None, 8, 8, 25)	0	batch_normalization_83[0][0]
conv2d_83 (Conv2D)	(None, 8, 8, 5)	3125	activation_82[0][0]
concatenate_76 (Concatenate)	(None, 8, 8, 30)	0	concatenate_75[0][0] conv2d_83[0][0]
batch_normalization_84 (BatchNo	(None, 8, 8, 30)	120	concatenate_76[0][0]
activation_83 (Activation)	(None, 8, 8, 30)	0	batch_normalization_84[0][0]

conv2d_84 (Conv2D)	(None, 8, 8, 5)	3750	activation_83[0][0]
concatenate_77 (Concatenate)	(None, 8, 8, 35)	0	concatenate_76[0][0] conv2d_84[0][0]
batch_normalization_85 (Batch Normalization)	(None, 8, 8, 35)	140	concatenate_77[0][0]
activation_84 (Activation)	(None, 8, 8, 35)	0	batch_normalization_85[0][0]
conv2d_85 (Conv2D)	(None, 8, 8, 5)	4375	activation_84[0][0]
concatenate_78 (Concatenate)	(None, 8, 8, 40)	0	concatenate_77[0][0] conv2d_85[0][0]
batch_normalization_86 (Batch Normalization)	(None, 8, 8, 40)	160	concatenate_78[0][0]
activation_85 (Activation)	(None, 8, 8, 40)	0	batch_normalization_86[0][0]
conv2d_86 (Conv2D)	(None, 8, 8, 5)	5000	activation_85[0][0]
concatenate_79 (Concatenate)	(None, 8, 8, 45)	0	concatenate_78[0][0] conv2d_86[0][0]
batch_normalization_87 (Batch Normalization)	(None, 8, 8, 45)	180	concatenate_79[0][0]
activation_86 (Activation)	(None, 8, 8, 45)	0	batch_normalization_87[0][0]
conv2d_87 (Conv2D)	(None, 8, 8, 5)	5625	activation_86[0][0]
concatenate_80 (Concatenate)	(None, 8, 8, 50)	0	concatenate_79[0][0] conv2d_87[0][0]
batch_normalization_88 (Batch Normalization)	(None, 8, 8, 50)	200	concatenate_80[0][0]
activation_87 (Activation)	(None, 8, 8, 50)	0	batch_normalization_88[0][0]
conv2d_88 (Conv2D)	(None, 8, 8, 5)	6250	activation_87[0][0]
concatenate_81 (Concatenate)	(None, 8, 8, 55)	0	concatenate_80[0][0] conv2d_88[0][0]
batch_normalization_89 (Batch Normalization)	(None, 8, 8, 55)	220	concatenate_81[0][0]
activation_88 (Activation)	(None, 8, 8, 55)	0	batch_normalization_89[0][0]
conv2d_89 (Conv2D)	(None, 8, 8, 5)	6875	activation_88[0][0]
concatenate_82 (Concatenate)	(None, 8, 8, 60)	0	concatenate_81[0][0] conv2d_89[0][0]
batch_normalization_90 (Batch Normalization)	(None, 8, 8, 60)	240	concatenate_82[0][0]
activation_89 (Activation)	(None, 8, 8, 60)	0	batch_normalization_90[0][0]
conv2d_90 (Conv2D)	(None, 8, 8, 5)	7500	activation_89[0][0]
concatenate_83 (Concatenate)	(None, 8, 8, 65)	0	concatenate_82[0][0] conv2d_90[0][0]

batch_normalization_91 (BatchNo	(None, 8, 8, 65)	260	concatenate_83[0][0]
activation_90 (Activation)	(None, 8, 8, 65)	0	batch_normalization_91[0][0]
conv2d_91 (Conv2D)	(None, 8, 8, 5)	8125	activation_90[0][0]
average_pooling2d_6 (AveragePoo	(None, 4, 4, 5)	0	conv2d_91[0][0]
batch_normalization_92 (BatchNo	(None, 4, 4, 5)	20	average_pooling2d_6[0][0]
activation_91 (Activation)	(None, 4, 4, 5)	0	batch_normalization_92[0][0]
conv2d_92 (Conv2D)	(None, 4, 4, 5)	625	activation_91[0][0]
concatenate_84 (Concatenate)	(None, 4, 4, 10)	0	average_pooling2d_6[0][0] conv2d_92[0][0]
batch_normalization_93 (BatchNo	(None, 4, 4, 10)	40	concatenate_84[0][0]
activation_92 (Activation)	(None, 4, 4, 10)	0	batch_normalization_93[0][0]
conv2d_93 (Conv2D)	(None, 4, 4, 5)	1250	activation_92[0][0]
concatenate_85 (Concatenate)	(None, 4, 4, 15)	0	concatenate_84[0][0] conv2d_93[0][0]
batch_normalization_94 (BatchNo	(None, 4, 4, 15)	60	concatenate_85[0][0]
activation_93 (Activation)	(None, 4, 4, 15)	0	batch_normalization_94[0][0]
conv2d_94 (Conv2D)	(None, 4, 4, 5)	1875	activation_93[0][0]
concatenate_86 (Concatenate)	(None, 4, 4, 20)	0	concatenate_85[0][0] conv2d_94[0][0]
batch_normalization_95 (BatchNo	(None, 4, 4, 20)	80	concatenate_86[0][0]
activation_94 (Activation)	(None, 4, 4, 20)	0	batch_normalization_95[0][0]
conv2d_95 (Conv2D)	(None, 4, 4, 5)	2500	activation_94[0][0]
concatenate_87 (Concatenate)	(None, 4, 4, 25)	0	concatenate_86[0][0] conv2d_95[0][0]
batch_normalization_96 (BatchNo	(None, 4, 4, 25)	100	concatenate_87[0][0]
activation_95 (Activation)	(None, 4, 4, 25)	0	batch_normalization_96[0][0]
conv2d_96 (Conv2D)	(None, 4, 4, 5)	3125	activation_95[0][0]
concatenate_88 (Concatenate)	(None, 4, 4, 30)	0	concatenate_87[0][0] conv2d_96[0][0]
batch_normalization_97 (BatchNo	(None, 4, 4, 30)	120	concatenate_88[0][0]
activation_96 (Activation)	(None, 4, 4, 30)	0	batch_normalization_97[0][0]
conv2d_97 (Conv2D)	(None, 4, 4, 5)	3750	activation_96[0][0]

concatenate_89 (Concatenate)	(None, 4, 4, 35)	0	concatenate_88[0][0] conv2d_97[0][0]
batch_normalization_98 (BatchNo	(None, 4, 4, 35)	140	concatenate_89[0][0]
activation_97 (Activation)	(None, 4, 4, 35)	0	batch_normalization_98[0][0]
conv2d_98 (Conv2D)	(None, 4, 4, 5)	4375	activation_97[0][0]
concatenate_90 (Concatenate)	(None, 4, 4, 40)	0	concatenate_89[0][0] conv2d_98[0][0]
batch_normalization_99 (BatchNo	(None, 4, 4, 40)	160	concatenate_90[0][0]
activation_98 (Activation)	(None, 4, 4, 40)	0	batch_normalization_99[0][0]
conv2d_99 (Conv2D)	(None, 4, 4, 5)	5000	activation_98[0][0]
concatenate_91 (Concatenate)	(None, 4, 4, 45)	0	concatenate_90[0][0] conv2d_99[0][0]
batch_normalization_100 (BatchN	(None, 4, 4, 45)	180	concatenate_91[0][0]
activation_99 (Activation)	(None, 4, 4, 45)	0	batch_normalization_100[0][0]
conv2d_100 (Conv2D)	(None, 4, 4, 5)	5625	activation_99[0][0]
concatenate_92 (Concatenate)	(None, 4, 4, 50)	0	concatenate_91[0][0] conv2d_100[0][0]
batch_normalization_101 (BatchN	(None, 4, 4, 50)	200	concatenate_92[0][0]
activation_100 (Activation)	(None, 4, 4, 50)	0	batch_normalization_101[0][0]
conv2d_101 (Conv2D)	(None, 4, 4, 5)	6250	activation_100[0][0]
concatenate_93 (Concatenate)	(None, 4, 4, 55)	0	concatenate_92[0][0] conv2d_101[0][0]
batch_normalization_102 (BatchN	(None, 4, 4, 55)	220	concatenate_93[0][0]
activation_101 (Activation)	(None, 4, 4, 55)	0	batch_normalization_102[0][0]
conv2d_102 (Conv2D)	(None, 4, 4, 5)	6875	activation_101[0][0]
concatenate_94 (Concatenate)	(None, 4, 4, 60)	0	concatenate_93[0][0] conv2d_102[0][0]
batch_normalization_103 (BatchN	(None, 4, 4, 60)	240	concatenate_94[0][0]
activation_102 (Activation)	(None, 4, 4, 60)	0	batch_normalization_103[0][0]
conv2d_103 (Conv2D)	(None, 4, 4, 5)	7500	activation_102[0][0]
concatenate_95 (Concatenate)	(None, 4, 4, 65)	0	concatenate_94[0][0] conv2d_103[0][0]
batch_normalization_104 (BatchN	(None, 4, 4, 65)	260	concatenate_95[0][0]



activation_103 (Activation)	(None, 4, 4, 65)	0	batch_normalization_104[0][0]
average_pooling2d_7 (AveragePool)	(None, 2, 2, 65)	0	activation_103[0][0]
flatten_1 (Flatten)	(None, 260)	0	average_pooling2d_7[0][0]
dense_1 (Dense)	(None, 10)	2610	flatten_1[0][0]
=====			
Total params: 746,808			
Trainable params: 740,834			
Non-trainable params: 5,974			
=====			

**we have 0.8 million trainable params**

```
In [24]: 1 from keras.preprocessing.image import ImageDataGenerator
2
3 ###Data Augmentation
4 ### https://nanonets.com/blog/data-augmentation-how-to-use-deep-learning-when-you-have-limited-data-part-2/
5 # rotation_range = 15, horizontal_flip = True, width_shift_range = 0.1, height_shift_range = 0.1, zoom_range = 0.2, shear_range = 15
6 def augment_data():
7     datagen = ImageDataGenerator(
8         rotation_range=15,
9         width_shift_range=0.1,
10        height_shift_range=0.09,
11        horizontal_flip=True,
12        zoom_range=0.2,
13        shear_range=15
14    )
15    return datagen
16 datagen = augment_data()
17 datagen.fit(X_train)
```

```
In [ ]: 1 ### https://keras.io/api/callbacks/model_checkpoint/
```

```
In [ ]: 1 # from google.colab import drive
2 # drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: 1 path = 'best_weights_epoch20.h5'
2 checkpoint = keras.callbacks.ModelCheckpoint(
3     filepath=path,
4     monitor='val_accuracy',
5     mode='max',
6     save_best_only=True)
```

```
In [ ]: 1 # from google.colab import drive
2 # drive.mount('/content/drive')
```

```
In [31]: 1 batch_size=64
2 model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
3 history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=batch_size),\
4                             steps_per_epoch=X_train.shape[0] // batch_size, epochs=20,\
5                             verbose=1, validation_data=(X_test, y_test), callbacks=[checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

Epoch 1/20

781/781 [=====] - 153s 148ms/step - loss: 1.7250 - accuracy: 0.3659 - val\_loss: 1.6352 - val\_accuracy: 0.4372

Epoch 2/20

781/781 [=====] - 114s 145ms/step - loss: 1.4152 - accuracy: 0.4820 - val\_loss: 1.3549 - val\_accuracy: 0.5060

Epoch 3/20

781/781 [=====] - 116s 148ms/step - loss: 1.2497 - accuracy: 0.5505 - val\_loss: 1.5288 - val\_accuracy: 0.4906

Epoch 4/20

781/781 [=====] - 119s 153ms/step - loss: 1.1252 - accuracy: 0.5960 - val\_loss: 1.1729 - val\_accuracy: 0.6021

Epoch 5/20

781/781 [=====] - 116s 148ms/step - loss: 1.0303 - accuracy: 0.6335 - val\_loss: 1.1565 - val\_accuracy: 0.6027

Epoch 6/20

781/781 [=====] - 116s 149ms/step - loss: 0.9658 - accuracy: 0.6551 - val\_loss: 1.0424 - val\_accuracy: 0.6391

Epoch 7/20

781/781 [=====] - 116s 149ms/step - loss: 0.9110 - accuracy: 0.6756 - val\_loss: 0.8783 - val\_accuracy: 0.6937

Epoch 8/20

781/781 [=====] - 116s 148ms/step - loss: 0.8595 - accuracy: 0.6953 - val\_loss: 0.9983 - val\_accuracy: 0.6614

Epoch 9/20

781/781 [=====] - 115s 147ms/step - loss: 0.8199 - accuracy: 0.7094 - val\_loss: 1.0162 - val\_accuracy: 0.6650

Epoch 10/20

781/781 [=====] - 115s 147ms/step - loss: 0.7833 - accuracy: 0.7249 - val\_loss: 1.1882 - val\_accuracy: 0.6387

Epoch 11/20

781/781 [=====] - 116s 149ms/step - loss: 0.7541 - accuracy: 0.7358 - val\_loss: 0.9855 - val\_accuracy: 0.6864

Epoch 12/20

781/781 [=====] - 116s 149ms/step - loss: 0.7270 - accuracy: 0.7451 - val\_loss: 0.8677 - val\_accuracy: 0.7156

Epoch 13/20

781/781 [=====] - 117s 149ms/step - loss: 0.6991 - accuracy: 0.7536 - val\_loss: 0.6496 - val\_accuracy: 0.7756

Epoch 14/20

781/781 [=====] - 115s 147ms/step - loss: 0.6761 - accuracy: 0.7640 - val\_loss: 0.7735 - val\_accuracy: 0.7373

Epoch 15/20

781/781 [=====] - 115s 147ms/step - loss: 0.6579 - accuracy: 0.7697 - val\_loss: 0.6855 - val\_accuracy: 0.7683

Epoch 16/20

781/781 [=====] - 119s 152ms/step - loss: 0.6426 - accuracy: 0.7749 - val\_loss: 0.8081 - val\_accuracy: 0.7320

Epoch 17/20

781/781 [=====] - 116s 149ms/step - loss: 0.6255 - accuracy: 0.7820 - val\_loss: 0.6836 - val\_accuracy: 0.7732

Epoch 18/20

781/781 [=====] - 118s 151ms/step - loss: 0.6110 - accuracy: 0.7869 - val\_loss: 0.6480 - val\_accuracy: 0.7792

Epoch 19/20

781/781 [=====] - 115s 148ms/step - loss: 0.5909 - accuracy: 0.7941 - val\_loss: 0.5855 - val\_accuracy: 0.8045

Epoch 20/20

781/781 [=====] - 116s 148ms/step - loss: 0.5787 - accuracy: 0.7957 - val\_loss: 0.6022 - val\_accuracy: 0.7945

```
In [ ]: ▶ 1 # Test the model
          2 model.load_weights(path)
          3 score = model.evaluate(X_test, y_test, verbose=1)
          4 print('Test loss:', score[0])
          5 print('Test accuracy:', score[1])
```

```
313/313 [=====] - 20s 65ms/step - loss: 0.6417 - accuracy: 0.7849
Test loss: 0.641657829284668
Test accuracy: 0.7849000096321106
```

```
In [ ]: ▶ 1 path = 'best_weights_epoch55.h5'
```

```
In [ ]: 1 model.load_weights('best_weights_epoch40.h5')
2         checkpoint = keras.callbacks.ModelCheckpoint(
3             filepath=path,
4             monitor='val_accuracy',
5             mode='max',
6             save_best_only=True)
7         model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
8         history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=batch_size),\
9                                     steps_per_epoch=X_train.shape[0] // batch_size, epochs=15,\
10                                    verbose=1, validation_data=(X_test, y_test), callbacks=[checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

```
Epoch 1/15
781/781 [=====] - 156s 151ms/step - loss: 0.4178 - accuracy: 0.8550 - val_loss: 0.5698 - val_accuracy: 0.8134
Epoch 2/15
781/781 [=====] - 115s 147ms/step - loss: 0.4090 - accuracy: 0.8568 - val_loss: 0.5355 - val_accuracy: 0.8202
Epoch 3/15
781/781 [=====] - 115s 147ms/step - loss: 0.4031 - accuracy: 0.8618 - val_loss: 0.5208 - val_accuracy: 0.8301
Epoch 4/15
781/781 [=====] - 115s 147ms/step - loss: 0.3976 - accuracy: 0.8615 - val_loss: 0.5456 - val_accuracy: 0.8216
Epoch 5/15
781/781 [=====] - 115s 148ms/step - loss: 0.3936 - accuracy: 0.8625 - val_loss: 0.4491 - val_accuracy: 0.8510
Epoch 6/15
781/781 [=====] - 115s 148ms/step - loss: 0.3928 - accuracy: 0.8623 - val_loss: 0.5136 - val_accuracy: 0.8268
Epoch 7/15
781/781 [=====] - 115s 147ms/step - loss: 0.3825 - accuracy: 0.8661 - val_loss: 0.4635 - val_accuracy: 0.8492
Epoch 8/15
781/781 [=====] - 115s 147ms/step - loss: 0.3789 - accuracy: 0.8672 - val_loss: 0.5166 - val_accuracy: 0.8311
Epoch 9/15
781/781 [=====] - 115s 147ms/step - loss: 0.3761 - accuracy: 0.8695 - val_loss: 0.5255 - val_accuracy: 0.8289
Epoch 10/15
781/781 [=====] - 118s 151ms/step - loss: 0.3741 - accuracy: 0.8690 - val_loss: 0.5462 - val_accuracy: 0.8299
Epoch 11/15
781/781 [=====] - 115s 148ms/step - loss: 0.3642 - accuracy: 0.8716 - val_loss: 0.4804 - val_accuracy: 0.8439
Epoch 12/15
781/781 [=====] - 117s 150ms/step - loss: 0.3646 - accuracy: 0.8729 - val_loss: 0.4253 - val_accuracy: 0.8579
Epoch 13/15
781/781 [=====] - 115s 147ms/step - loss: 0.3571 - accuracy: 0.8751 - val_loss: 0.6388 - val_accuracy: 0.8069
Epoch 14/15
781/781 [=====] - 115s 147ms/step - loss: 0.3539 - accuracy: 0.8780 - val_loss: 0.5496 - val_accuracy: 0.8301
Epoch 15/15
781/781 [=====] - 115s 147ms/step - loss: 0.3511 - accuracy: 0.8766 - val_loss: 0.4427 - val_accuracy: 0.8541
```

```
In [ ]: 1 !cp -r '/content/best_weights_epoch55.h5' '/content/drive/MyDrive/'
```

```
In [ ]: 1 path = 'best_weights_epoch85.h5'
2         checkpoint = keras.callbacks.ModelCheckpoint(
3             filepath=path,
4             monitor='val_accuracy',
5             mode='max',
6             save_best_only=True)
```

```
In [ ]: 1 model.load_weights('best_weights_epoch55.h5')
2 model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
3 history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=batch_size),\
4                             steps_per_epoch=X_train.shape[0] // batch_size, epochs=30,\
5                             verbose=1, validation_data=(X_test, y_test), callbacks=[checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

```
Epoch 1/30
781/781 [=====] - 157s 153ms/step - loss: 0.3638 - accuracy: 0.8736 - val_loss: 0.4842 - val_accuracy: 0.8436
Epoch 2/30
781/781 [=====] - 115s 147ms/step - loss: 0.3564 - accuracy: 0.8768 - val_loss: 0.6951 - val_accuracy: 0.7971
Epoch 3/30
781/781 [=====] - 118s 151ms/step - loss: 0.3524 - accuracy: 0.8761 - val_loss: 0.5342 - val_accuracy: 0.8328
Epoch 4/30
781/781 [=====] - 115s 148ms/step - loss: 0.3473 - accuracy: 0.8776 - val_loss: 0.4074 - val_accuracy: 0.8682
Epoch 5/30
781/781 [=====] - 115s 147ms/step - loss: 0.3444 - accuracy: 0.8798 - val_loss: 0.4847 - val_accuracy: 0.8424
Epoch 6/30
781/781 [=====] - 117s 150ms/step - loss: 0.3414 - accuracy: 0.8816 - val_loss: 0.4684 - val_accuracy: 0.8437
Epoch 7/30
781/781 [=====] - 117s 150ms/step - loss: 0.3393 - accuracy: 0.8811 - val_loss: 0.4696 - val_accuracy: 0.8452
Epoch 8/30
781/781 [=====] - 115s 147ms/step - loss: 0.3337 - accuracy: 0.8816 - val_loss: 0.4637 - val_accuracy: 0.8458
Epoch 9/30
781/781 [=====] - 114s 146ms/step - loss: 0.3331 - accuracy: 0.8840 - val_loss: 0.4954 - val_accuracy: 0.8412
Epoch 10/30
781/781 [=====] - 115s 147ms/step - loss: 0.3297 - accuracy: 0.8843 - val_loss: 0.5015 - val_accuracy: 0.8437
Epoch 11/30
781/781 [=====] - 115s 147ms/step - loss: 0.3295 - accuracy: 0.8850 - val_loss: 0.5027 - val_accuracy: 0.8419
Epoch 12/30
781/781 [=====] - 115s 147ms/step - loss: 0.3216 - accuracy: 0.8875 - val_loss: 0.4757 - val_accuracy: 0.8480
Epoch 13/30
781/781 [=====] - 115s 147ms/step - loss: 0.3182 - accuracy: 0.8879 - val_loss: 0.5190 - val_accuracy: 0.8341
Epoch 14/30
781/781 [=====] - 115s 147ms/step - loss: 0.3188 - accuracy: 0.8900 - val_loss: 0.5112 - val_accuracy: 0.8456
Epoch 15/30
781/781 [=====] - 114s 146ms/step - loss: 0.3218 - accuracy: 0.8868 - val_loss: 0.3919 - val_accuracy: 0.8691
Epoch 16/30
781/781 [=====] - 114s 146ms/step - loss: 0.3119 - accuracy: 0.8889 - val_loss: 0.3960 - val_accuracy: 0.8684
Epoch 17/30
781/781 [=====] - 115s 147ms/step - loss: 0.3085 - accuracy: 0.8924 - val_loss: 0.5192 - val_accuracy: 0.8417
Epoch 18/30
781/781 [=====] - 115s 147ms/step - loss: 0.3053 - accuracy: 0.8922 - val_loss: 0.5063 - val_accuracy: 0.8392
Epoch 19/30
781/781 [=====] - 114s 146ms/step - loss: 0.3090 - accuracy: 0.8928 - val_loss: 0.4406 - val_accuracy: 0.8595
Epoch 20/30
781/781 [=====] - 115s 147ms/step - loss: 0.3005 - accuracy: 0.8957 - val_loss: 0.4997 - val_accuracy: 0.8426
Epoch 21/30
781/781 [=====] - 115s 147ms/step - loss: 0.2984 - accuracy: 0.8957 - val_loss: 0.4546 - val_accuracy: 0.8527
Epoch 22/30
781/781 [=====] - 115s 147ms/step - loss: 0.2953 - accuracy: 0.8966 - val_loss: 0.6197 - val_accuracy: 0.8173
Epoch 23/30
781/781 [=====] - 114s 146ms/step - loss: 0.2953 - accuracy: 0.8970 - val_loss: 0.5026 - val_accuracy: 0.8380
Epoch 24/30
```

```

781/781 [=====] - 115s 147ms/step - loss: 0.2935 - accuracy: 0.8977 - val_loss: 0.4265 - val_accuracy: 0.8623
Epoch 25/30
781/781 [=====] - 115s 148ms/step - loss: 0.2885 - accuracy: 0.8975 - val_loss: 0.4517 - val_accuracy: 0.8579
Epoch 26/30
781/781 [=====] - 116s 148ms/step - loss: 0.2894 - accuracy: 0.8981 - val_loss: 0.4063 - val_accuracy: 0.8677
Epoch 27/30
781/781 [=====] - 115s 147ms/step - loss: 0.2833 - accuracy: 0.8997 - val_loss: 0.5102 - val_accuracy: 0.8440
Epoch 28/30
781/781 [=====] - 118s 151ms/step - loss: 0.2824 - accuracy: 0.9003 - val_loss: 0.4280 - val_accuracy: 0.8647
Epoch 29/30
781/781 [=====] - 115s 148ms/step - loss: 0.2812 - accuracy: 0.9013 - val_loss: 0.3993 - val_accuracy: 0.8751
Epoch 30/30
781/781 [=====] - 115s 148ms/step - loss: 0.2799 - accuracy: 0.9020 - val_loss: 0.4679 - val_accuracy: 0.8547

```

```
In [ ]: 1 !cp -r '/content/best_weights_epoch85.h5' '/content/drive/MyDrive/'
```

```
In [ ]: 1 path = 'best_weights_epoch115.h5'
2 checkpoint = keras.callbacks.ModelCheckpoint(
3     filepath=path,
4     monitor='val_accuracy',
5     mode='max',
6     save_best_only=True)
```

```
In [ ]: 1 model.load_weights('best_weights_epoch85.h5')
2 model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
3 history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=batch_size),\
4                             steps_per_epoch=X_train.shape[0] // batch_size, epochs=30,\
5                             verbose=1, validation_data=(X_test, y_test), callbacks=[checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

Epoch 1/30

781/781 [=====] - 159s 155ms/step - loss: 0.2775 - accuracy: 0.9017 - val\_loss: 0.4250 - val\_accuracy: 0.8604

Epoch 2/30

781/781 [=====] - 121s 154ms/step - loss: 0.2778 - accuracy: 0.9017 - val\_loss: 0.3903 - val\_accuracy: 0.8730

Epoch 3/30

781/781 [=====] - 120s 154ms/step - loss: 0.2780 - accuracy: 0.9028 - val\_loss: 0.4269 - val\_accuracy: 0.8666

Epoch 4/30

781/781 [=====] - 117s 150ms/step - loss: 0.2742 - accuracy: 0.9043 - val\_loss: 0.4757 - val\_accuracy: 0.8550

Epoch 5/30

781/781 [=====] - 117s 150ms/step - loss: 0.2723 - accuracy: 0.9037 - val\_loss: 0.4582 - val\_accuracy: 0.8588

Epoch 6/30

781/781 [=====] - 118s 151ms/step - loss: 0.2720 - accuracy: 0.9042 - val\_loss: 0.4763 - val\_accuracy: 0.8529

Epoch 7/30

781/781 [=====] - 117s 150ms/step - loss: 0.2643 - accuracy: 0.9076 - val\_loss: 0.4387 - val\_accuracy: 0.8646

Epoch 8/30

781/781 [=====] - 117s 150ms/step - loss: 0.2738 - accuracy: 0.9043 - val\_loss: 0.5215 - val\_accuracy: 0.8448

Epoch 9/30

781/781 [=====] - 117s 150ms/step - loss: 0.2650 - accuracy: 0.9074 - val\_loss: 0.4282 - val\_accuracy: 0.8688

Epoch 10/30

781/781 [=====] - 118s 151ms/step - loss: 0.2647 - accuracy: 0.9085 - val\_loss: 0.4375 - val\_accuracy: 0.8617

Epoch 11/30

781/781 [=====] - 118s 151ms/step - loss: 0.2594 - accuracy: 0.9094 - val\_loss: 0.5873 - val\_accuracy: 0.8325

Epoch 12/30

781/781 [=====] - 120s 154ms/step - loss: 0.2656 - accuracy: 0.9068 - val\_loss: 0.4245 - val\_accuracy: 0.8668

Epoch 13/30

781/781 [=====] - 118s 151ms/step - loss: 0.2659 - accuracy: 0.9062 - val\_loss: 0.4918 - val\_accuracy: 0.8467

Epoch 14/30

781/781 [=====] - 117s 150ms/step - loss: 0.2585 - accuracy: 0.9099 - val\_loss: 0.3979 - val\_accuracy: 0.8731

Epoch 15/30

781/781 [=====] - 120s 154ms/step - loss: 0.2574 - accuracy: 0.9093 - val\_loss: 0.4473 - val\_accuracy: 0.8671

Epoch 16/30

781/781 [=====] - 118s 151ms/step - loss: 0.2533 - accuracy: 0.9107 - val\_loss: 0.3749 - val\_accuracy: 0.8799

Epoch 17/30

781/781 [=====] - 118s 151ms/step - loss: 0.2496 - accuracy: 0.9116 - val\_loss: 0.4054 - val\_accuracy: 0.8706

Epoch 18/30

781/781 [=====] - 118s 151ms/step - loss: 0.2510 - accuracy: 0.9122 - val\_loss: 0.4819 - val\_accuracy: 0.8544

Epoch 19/30

781/781 [=====] - 117s 150ms/step - loss: 0.2489 - accuracy: 0.9118 - val\_loss: 0.4347 - val\_accuracy: 0.8655

Epoch 20/30

781/781 [=====] - 120s 154ms/step - loss: 0.2465 - accuracy: 0.9130 - val\_loss: 0.4244 - val\_accuracy: 0.8721

Epoch 21/30

781/781 [=====] - 117s 150ms/step - loss: 0.2492 - accuracy: 0.9132 - val\_loss: 0.4077 - val\_accuracy: 0.8744

Epoch 22/30

781/781 [=====] - 118s 151ms/step - loss: 0.2456 - accuracy: 0.9139 - val\_loss: 0.3849 - val\_accuracy: 0.8797

Epoch 23/30

781/781 [=====] - 118s 151ms/step - loss: 0.2471 - accuracy: 0.9132 - val\_loss: 0.4045 - val\_accuracy: 0.8753

Epoch 24/30

781/781 [=====] - 118s 151ms/step - loss: 0.2456 - accuracy: 0.9129 - val\_loss: 0.3963 - val\_accuracy: 0.8789

```
Epoch 25/30
781/781 [=====] - 118s 151ms/step - loss: 0.2396 - accuracy: 0.9155 - val_loss: 0.4283 - val_accuracy: 0.8662
Epoch 26/30
781/781 [=====] - 121s 154ms/step - loss: 0.2410 - accuracy: 0.9154 - val_loss: 0.4251 - val_accuracy: 0.8658
Epoch 27/30
781/781 [=====] - 117s 149ms/step - loss: 0.2362 - accuracy: 0.9173 - val_loss: 0.4279 - val_accuracy: 0.8692
Epoch 28/30
781/781 [=====] - 118s 150ms/step - loss: 0.2357 - accuracy: 0.9173 - val_loss: 0.4325 - val_accuracy: 0.8688
Epoch 29/30
781/781 [=====] - 118s 150ms/step - loss: 0.2387 - accuracy: 0.9164 - val_loss: 0.4364 - val_accuracy: 0.8709
Epoch 30/30
781/781 [=====] - 120s 154ms/step - loss: 0.2322 - accuracy: 0.9190 - val_loss: 0.3871 - val_accuracy: 0.8788
```

```
In [ ]: 1 ##### visualize loss and layers for 100 epochs
2 import seaborn as sns
3 def plot_accuracy(train_loss,test_loss):
4     fig = plt.figure(figsize=(9,5))
5     sns.lineplot(np.arange(100),train_loss,label='train_loss')
6     sns.lineplot(np.arange(100),test_loss,label='test_loss')
7     plt.title('Test-Train Accuracy')
8     plt.xlabel('epochs')
9     plt.ylabel('loss')
10    plt.show()
11    plot_accuracy(history['accuracy'][:100],history['val_accuracy'][:100])
```



```
In [ ]: 1 path = 'best_weights_epoch135.h5'
2 checkpoint = keras.callbacks.ModelCheckpoint(
3     filepath=path,
4     monitor='val_accuracy',
5     mode='max',
6     save_best_only=True)
```



```
In [ ]: 1 model.load_weights('best_weights_epoch115.h5')
2 model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
3 history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=batch_size),\
4                             steps_per_epoch=X_train.shape[0] // batch_size, epochs=20,\
5                             verbose=1, validation_data=(X_test, y_test), callbacks=[checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

Epoch 1/20

781/781 [=====] - 123s 152ms/step - loss: 0.2464 - accuracy: 0.9125 - val\_loss: 0.4448 - val\_accuracy: 0.8660

Epoch 2/20

781/781 [=====] - 117s 150ms/step - loss: 0.2394 - accuracy: 0.9154 - val\_loss: 0.4264 - val\_accuracy: 0.8678

Epoch 3/20

781/781 [=====] - 117s 150ms/step - loss: 0.2428 - accuracy: 0.9146 - val\_loss: 0.4666 - val\_accuracy: 0.8594

Epoch 4/20

781/781 [=====] - 117s 150ms/step - loss: 0.2419 - accuracy: 0.9137 - val\_loss: 0.4198 - val\_accuracy: 0.8711

Epoch 5/20

781/781 [=====] - 118s 151ms/step - loss: 0.2370 - accuracy: 0.9173 - val\_loss: 0.4397 - val\_accuracy: 0.8677

Epoch 6/20

781/781 [=====] - 118s 151ms/step - loss: 0.2396 - accuracy: 0.9161 - val\_loss: 0.4658 - val\_accuracy: 0.8606

Epoch 7/20

781/781 [=====] - 117s 150ms/step - loss: 0.2361 - accuracy: 0.9169 - val\_loss: 0.4918 - val\_accuracy: 0.8523

Epoch 8/20

781/781 [=====] - 120s 154ms/step - loss: 0.2391 - accuracy: 0.9155 - val\_loss: 0.3854 - val\_accuracy: 0.8811

Epoch 9/20

781/781 [=====] - 121s 154ms/step - loss: 0.2364 - accuracy: 0.9173 - val\_loss: 0.3571 - val\_accuracy: 0.8845

Epoch 10/20

781/781 [=====] - 118s 151ms/step - loss: 0.2327 - accuracy: 0.9168 - val\_loss: 0.4488 - val\_accuracy: 0.8686

Epoch 11/20

781/781 [=====] - 118s 150ms/step - loss: 0.2303 - accuracy: 0.9194 - val\_loss: 0.4265 - val\_accuracy: 0.8689

Epoch 12/20

781/781 [=====] - 121s 154ms/step - loss: 0.2298 - accuracy: 0.9200 - val\_loss: 0.4837 - val\_accuracy: 0.8580

Epoch 13/20

781/781 [=====] - 117s 150ms/step - loss: 0.2275 - accuracy: 0.9199 - val\_loss: 0.4764 - val\_accuracy: 0.8589

Epoch 14/20

781/781 [=====] - 120s 153ms/step - loss: 0.2247 - accuracy: 0.9207 - val\_loss: 0.4889 - val\_accuracy: 0.8556

Epoch 15/20

781/781 [=====] - 120s 153ms/step - loss: 0.2259 - accuracy: 0.9202 - val\_loss: 0.3972 - val\_accuracy: 0.8767

Epoch 16/20

781/781 [=====] - 117s 150ms/step - loss: 0.2285 - accuracy: 0.9189 - val\_loss: 0.4497 - val\_accuracy: 0.8646

Epoch 17/20

781/781 [=====] - 118s 150ms/step - loss: 0.2281 - accuracy: 0.9196 - val\_loss: 0.4538 - val\_accuracy: 0.8671

Epoch 18/20

781/781 [=====] - 117s 150ms/step - loss: 0.2272 - accuracy: 0.9210 - val\_loss: 0.4015 - val\_accuracy: 0.8732

Epoch 19/20

781/781 [=====] - 120s 154ms/step - loss: 0.2240 - accuracy: 0.9225 - val\_loss: 0.3444 - val\_accuracy: 0.8910

Epoch 20/20

781/781 [=====] - 118s 151ms/step - loss: 0.2217 - accuracy: 0.9205 - val\_loss: 0.4281 - val\_accuracy: 0.8700

```
In [ ]: ▶ 1 path = 'best_weights_epoch155.h5'
          2 checkpoint = keras.callbacks.ModelCheckpoint(
          3     filepath=path,
          4     monitor='val_accuracy',
          5     mode='max',
          6     save_best_only=True)
```

```
In [ ]: 1 model.load_weights('best_weights_epoch135.h5')
2 model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
3 history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=batch_size),\
4                             steps_per_epoch=X_train.shape[0] // batch_size, epochs=20,\
5                             verbose=1, validation_data=(X_test, y_test), callbacks=[checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and ")

```
Epoch 1/20
781/781 [=====] - 123s 152ms/step - loss: 0.2206 - accuracy: 0.9218 - val_loss: 0.3823 - val_accuracy: 0.8819
Epoch 2/20
781/781 [=====] - 120s 153ms/step - loss: 0.2188 - accuracy: 0.9247 - val_loss: 0.3703 - val_accuracy: 0.8842
Epoch 3/20
781/781 [=====] - 117s 150ms/step - loss: 0.2201 - accuracy: 0.9229 - val_loss: 0.4306 - val_accuracy: 0.8683
Epoch 4/20
781/781 [=====] - 118s 151ms/step - loss: 0.2195 - accuracy: 0.9221 - val_loss: 0.4011 - val_accuracy: 0.8799
Epoch 5/20
781/781 [=====] - 117s 150ms/step - loss: 0.2172 - accuracy: 0.9235 - val_loss: 0.5441 - val_accuracy: 0.8507
Epoch 6/20
781/781 [=====] - 118s 151ms/step - loss: 0.2171 - accuracy: 0.9238 - val_loss: 0.3852 - val_accuracy: 0.8819
Epoch 7/20
781/781 [=====] - 117s 150ms/step - loss: 0.2157 - accuracy: 0.9233 - val_loss: 0.3958 - val_accuracy: 0.8797
Epoch 8/20
781/781 [=====] - 118s 151ms/step - loss: 0.2126 - accuracy: 0.9238 - val_loss: 0.3593 - val_accuracy: 0.8885
Epoch 9/20
781/781 [=====] - 118s 151ms/step - loss: 0.2091 - accuracy: 0.9266 - val_loss: 0.4084 - val_accuracy: 0.8785
Epoch 10/20
781/781 [=====] - 118s 151ms/step - loss: 0.2155 - accuracy: 0.9242 - val_loss: 0.3320 - val_accuracy: 0.8986
Epoch 11/20
781/781 [=====] - 118s 151ms/step - loss: 0.2121 - accuracy: 0.9251 - val_loss: 0.4551 - val_accuracy: 0.8667
Epoch 12/20
781/781 [=====] - 117s 149ms/step - loss: 0.2092 - accuracy: 0.9267 - val_loss: 0.3868 - val_accuracy: 0.8774
Epoch 13/20
781/781 [=====] - 116s 149ms/step - loss: 0.2101 - accuracy: 0.9244 - val_loss: 0.4343 - val_accuracy: 0.8744
Epoch 14/20
781/781 [=====] - 120s 153ms/step - loss: 0.2075 - accuracy: 0.9261 - val_loss: 0.5316 - val_accuracy: 0.8532
Epoch 15/20
781/781 [=====] - 118s 151ms/step - loss: 0.2021 - accuracy: 0.9279 - val_loss: 0.4122 - val_accuracy: 0.8747
Epoch 16/20
781/781 [=====] - 117s 150ms/step - loss: 0.2048 - accuracy: 0.9273 - val_loss: 0.4393 - val_accuracy: 0.8698
Epoch 17/20
781/781 [=====] - 120s 154ms/step - loss: 0.2031 - accuracy: 0.9291 - val_loss: 0.4312 - val_accuracy: 0.8744
Epoch 18/20
781/781 [=====] - 118s 150ms/step - loss: 0.2036 - accuracy: 0.9278 - val_loss: 0.3904 - val_accuracy: 0.8876
Epoch 19/20
781/781 [=====] - 118s 151ms/step - loss: 0.2002 - accuracy: 0.9280 - val_loss: 0.4365 - val_accuracy: 0.8758
Epoch 20/20
781/781 [=====] - 117s 150ms/step - loss: 0.2048 - accuracy: 0.9276 - val_loss: 0.4986 - val_accuracy: 0.8578
```

```
In [ ]: ▶ 1 path = 'best_weights_epoch185.h5'
          2 checkpoint = keras.callbacks.ModelCheckpoint(
          3     filepath=path,
          4     monitor='val_accuracy',
          5     mode='max',
          6     save_best_only=True)
```

```
In [ ]: 1 model.load_weights('best_weights_epoch155.h5')
2 model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])
3 history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=batch_size),\
4                             steps_per_epoch=X_train.shape[0] // batch_size, epochs=30,\
5                             verbose=1, validation_data=(X_test, y_test), callbacks=[checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

Epoch 1/30

781/781 [=====] - 289s 325ms/step - loss: 0.2166 - accuracy: 0.9234 - val\_loss: 0.3505 - val\_accuracy: 0.8853

Epoch 2/30

781/781 [=====] - 252s 322ms/step - loss: 0.2083 - accuracy: 0.9269 - val\_loss: 0.4370 - val\_accuracy: 0.8720

Epoch 3/30

781/781 [=====] - 252s 322ms/step - loss: 0.2087 - accuracy: 0.9259 - val\_loss: 0.3576 - val\_accuracy: 0.8885

Epoch 4/30

781/781 [=====] - 251s 321ms/step - loss: 0.2052 - accuracy: 0.9279 - val\_loss: 0.4450 - val\_accuracy: 0.8709

Epoch 5/30

781/781 [=====] - 251s 322ms/step - loss: 0.2071 - accuracy: 0.9259 - val\_loss: 0.3825 - val\_accuracy: 0.8850

Epoch 6/30

781/781 [=====] - 250s 320ms/step - loss: 0.2048 - accuracy: 0.9282 - val\_loss: 0.4575 - val\_accuracy: 0.8689

Epoch 7/30

781/781 [=====] - 251s 322ms/step - loss: 0.2066 - accuracy: 0.9261 - val\_loss: 0.4060 - val\_accuracy: 0.8728

Epoch 8/30

781/781 [=====] - 252s 322ms/step - loss: 0.2004 - accuracy: 0.9304 - val\_loss: 0.4743 - val\_accuracy: 0.8632

Epoch 9/30

781/781 [=====] - 252s 323ms/step - loss: 0.2040 - accuracy: 0.9277 - val\_loss: 0.5136 - val\_accuracy: 0.8573

Epoch 10/30

781/781 [=====] - 252s 323ms/step - loss: 0.2023 - accuracy: 0.9283 - val\_loss: 0.4570 - val\_accuracy: 0.8673

Epoch 11/30

781/781 [=====] - 252s 322ms/step - loss: 0.2016 - accuracy: 0.9285 - val\_loss: 0.3484 - val\_accuracy: 0.8919

Epoch 12/30

781/781 [=====] - 252s 322ms/step - loss: 0.2013 - accuracy: 0.9304 - val\_loss: 0.4652 - val\_accuracy: 0.8676

Epoch 13/30

781/781 [=====] - 251s 322ms/step - loss: 0.1974 - accuracy: 0.9303 - val\_loss: 0.4059 - val\_accuracy: 0.8796

Epoch 14/30

781/781 [=====] - 251s 322ms/step - loss: 0.1969 - accuracy: 0.9296 - val\_loss: 0.3836 - val\_accuracy: 0.8852

Epoch 15/30

781/781 [=====] - 252s 322ms/step - loss: 0.1967 - accuracy: 0.9305 - val\_loss: 0.4415 - val\_accuracy: 0.8747

Epoch 16/30

781/781 [=====] - 252s 323ms/step - loss: 0.1957 - accuracy: 0.9318 - val\_loss: 0.4962 - val\_accuracy: 0.8616

Epoch 17/30

781/781 [=====] - 252s 323ms/step - loss: 0.1955 - accuracy: 0.9305 - val\_loss: 0.4071 - val\_accuracy: 0.8809

Epoch 18/30

781/781 [=====] - 252s 322ms/step - loss: 0.1949 - accuracy: 0.9318 - val\_loss: 0.3903 - val\_accuracy: 0.8855

Epoch 19/30

781/781 [=====] - 252s 322ms/step - loss: 0.1969 - accuracy: 0.9311 - val\_loss: 0.5107 - val\_accuracy: 0.8593

Epoch 20/30

781/781 [=====] - 252s 322ms/step - loss: 0.1943 - accuracy: 0.9310 - val\_loss: 0.4738 - val\_accuracy: 0.8632

Epoch 21/30

781/781 [=====] - 252s 323ms/step - loss: 0.1913 - accuracy: 0.9332 - val\_loss: 0.4201 - val\_accuracy: 0.8777

Epoch 22/30

781/781 [=====] - 252s 323ms/step - loss: 0.1975 - accuracy: 0.9299 - val\_loss: 0.4235 - val\_accuracy: 0.8777

Epoch 23/30

781/781 [=====] - 252s 322ms/step - loss: 0.1936 - accuracy: 0.9306 - val\_loss: 0.4826 - val\_accuracy: 0.8654

Epoch 24/30

781/781 [=====] - 252s 323ms/step - loss: 0.1891 - accuracy: 0.9334 - val\_loss: 0.4679 - val\_accuracy: 0.8737

Epoch 25/30  
781/781 [=====] - 252s 323ms/step - loss: 0.1905 - accuracy: 0.9321 - val\_loss: 0.4436 - val\_accuracy: 0.8726  
Epoch 26/30  
781/781 [=====] - 251s 322ms/step - loss: 0.1893 - accuracy: 0.9331 - val\_loss: 0.4268 - val\_accuracy: 0.8731  
Epoch 27/30  
781/781 [=====] - 251s 322ms/step - loss: 0.1901 - accuracy: 0.9319 - val\_loss: 0.4381 - val\_accuracy: 0.8742  
Epoch 28/30  
781/781 [=====] - 252s 322ms/step - loss: 0.1898 - accuracy: 0.9334 - val\_loss: 0.4332 - val\_accuracy: 0.8778  
Epoch 29/30  
781/781 [=====] - 252s 322ms/step - loss: 0.1888 - accuracy: 0.9336 - val\_loss: 0.4167 - val\_accuracy: 0.8814  
Epoch 30/30  
781/781 [=====] - 250s 320ms/step - loss: 0.1852 - accuracy: 0.9344 - val\_loss: 0.4013 - val\_accuracy: 0.8836

let's use SGD optimizer with momentum 0.6 and learning rate 0.001 to converge faster to optimal solution

In [32]:

▶

1 from keras.optimizers import SGD

Type *Markdown* and LaTeX:  $\alpha^2$

In [ ]:

```

1
2 model.load_weights('best_weights_epoch185.h5')
3
4
5 path = 'best_weights_epoch205.h5'
6 checkpoint = keras.callbacks.ModelCheckpoint(
7     filepath=path,
8     monitor='val_accuracy',
9     mode='max',
10    save_best_only=True)
11
12 model.compile(loss='categorical_crossentropy', optimizer=SGD(momentum=0.6, learning_rate=0.001),
13               metrics=['accuracy'])
14 model.fit_generator(datagen.flow(X_train, y_train, batch_size),
15                    steps_per_epoch = X_train.shape[0]//batch_size,
16                    epochs = 20, validation_data =(X_test, y_test), callbacks = [checkpoint])

```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

Epoch 1/20

390/390 [=====] - 156s 297ms/step - loss: 0.1792 - accuracy: 0.9370 - val\_loss: 0.3536 - val\_accuracy: 0.8930

Epoch 2/20

390/390 [=====] - 113s 290ms/step - loss: 0.1709 - accuracy: 0.9398 - val\_loss: 0.3498 - val\_accuracy: 0.8952

Epoch 3/20

390/390 [=====] - 111s 284ms/step - loss: 0.1693 - accuracy: 0.9402 - val\_loss: 0.3473 - val\_accuracy: 0.8960

Epoch 4/20

390/390 [=====] - 111s 285ms/step - loss: 0.1648 - accuracy: 0.9417 - val\_loss: 0.3457 - val\_accuracy: 0.8964

Epoch 5/20

390/390 [=====] - 114s 291ms/step - loss: 0.1633 - accuracy: 0.9437 - val\_loss: 0.3446 - val\_accuracy: 0.8973

Epoch 6/20

390/390 [=====] - 111s 284ms/step - loss: 0.1607 - accuracy: 0.9442 - val\_loss: 0.3433 - val\_accuracy: 0.8976

Epoch 7/20

390/390 [=====] - 111s 283ms/step - loss: 0.1602 - accuracy: 0.9438 - val\_loss: 0.3416 - val\_accuracy: 0.8979

Epoch 8/20

390/390 [=====] - 110s 283ms/step - loss: 0.1598 - accuracy: 0.9446 - val\_loss: 0.3435 - val\_accuracy: 0.8974

Epoch 9/20

390/390 [=====] - 110s 282ms/step - loss: 0.1581 - accuracy: 0.9447 - val\_loss: 0.3417 - val\_accuracy: 0.8988

Epoch 10/20

390/390 [=====] - 110s 282ms/step - loss: 0.1569 - accuracy: 0.9449 - val\_loss: 0.3411 - val\_accuracy: 0.8988

Epoch 11/20

390/390 [=====] - 109s 280ms/step - loss: 0.1530 - accuracy: 0.9462 - val\_loss: 0.3409 - val\_accuracy: 0.8986

Epoch 12/20

390/390 [=====] - 110s 282ms/step - loss: 0.1561 - accuracy: 0.9458 - val\_loss: 0.3421 - val\_accuracy: 0.8988

Epoch 13/20

390/390 [=====] - 110s 283ms/step - loss: 0.1568 - accuracy: 0.9443 - val\_loss: 0.3422 - val\_accuracy: 0.8991

Epoch 14/20

390/390 [=====] - 109s 280ms/step - loss: 0.1578 - accuracy: 0.9441 - val\_loss: 0.3421 - val\_accuracy: 0.8994

Epoch 15/20

390/390 [=====] - 111s 283ms/step - loss: 0.1524 - accuracy: 0.9466 - val\_loss: 0.3418 - val\_accuracy: 0.8994

Epoch 16/20

390/390 [=====] - 110s 283ms/step - loss: 0.1519 - accuracy: 0.9466 - val\_loss: 0.3425 - val\_accuracy: 0.8995

Epoch 17/20

390/390 [=====] - 109s 280ms/step - loss: 0.1509 - accuracy: 0.9479 - val\_loss: 0.3415 - val\_accuracy: 0.8997

Epoch 18/20

390/390 [=====] - 110s 282ms/step - loss: 0.1516 - accuracy: 0.9476 - val\_loss: 0.3397 - val\_accuracy: 0.9006

Epoch 19/20

```
390/390 [=====] - 110s 283ms/step - loss: 0.1492 - accuracy: 0.9478 - val_loss: 0.3425 - val_accuracy: 0.8990
Epoch 20/20
390/390 [=====] - 109s 280ms/step - loss: 0.1478 - accuracy: 0.9485 - val_loss: 0.3391 - val_accuracy: 0.9001
```

Out[20]: <tensorflow.python.keras.callbacks.History at 0x7f857abd6d50>

```
In [ ]: 1 model.load_weights('best_weights_epoch205.h5')
        2
        3 path = 'best_weights_epoch215.h5'
        4 checkpoint = keras.callbacks.ModelCheckpoint(
        5     filepath=path,
        6     monitor='val_accuracy',
        7     mode='max',
        8     save_best_only=True)
        9
        10 model.compile(loss='categorical_crossentropy', optimizer=SGD(momentum=0.6, learning_rate=0.001),
        11                    metrics=['accuracy'])
        12 model.fit_generator(dagen.flow(X_train, y_train, batch_size),
        13                        steps_per_epoch = X_train.shape[0]//batch_size,
        14                        epochs = 10, validation_data =(X_test, y_test), callbacks = [checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

```
Epoch 1/10
390/390 [=====] - 110s 274ms/step - loss: 0.1534 - accuracy: 0.9461 - val_loss: 0.3416 - val_accuracy: 0.8999
Epoch 2/10
390/390 [=====] - 108s 278ms/step - loss: 0.1512 - accuracy: 0.9466 - val_loss: 0.3415 - val_accuracy: 0.9004
Epoch 3/10
390/390 [=====] - 112s 286ms/step - loss: 0.1474 - accuracy: 0.9488 - val_loss: 0.3395 - val_accuracy: 0.8999
Epoch 4/10
390/390 [=====] - 112s 287ms/step - loss: 0.1508 - accuracy: 0.9481 - val_loss: 0.3398 - val_accuracy: 0.9005
Epoch 5/10
390/390 [=====] - 110s 283ms/step - loss: 0.1464 - accuracy: 0.9482 - val_loss: 0.3400 - val_accuracy: 0.9002
Epoch 6/10
390/390 [=====] - 112s 288ms/step - loss: 0.1506 - accuracy: 0.9477 - val_loss: 0.3434 - val_accuracy: 0.8999
Epoch 7/10
390/390 [=====] - 109s 280ms/step - loss: 0.1479 - accuracy: 0.9475 - val_loss: 0.3406 - val_accuracy: 0.8999
Epoch 8/10
390/390 [=====] - 109s 280ms/step - loss: 0.1476 - accuracy: 0.9479 - val_loss: 0.3425 - val_accuracy: 0.8997
Epoch 9/10
390/390 [=====] - 109s 280ms/step - loss: 0.1484 - accuracy: 0.9467 - val_loss: 0.3421 - val_accuracy: 0.8995
Epoch 10/10
390/390 [=====] - 111s 284ms/step - loss: 0.1478 - accuracy: 0.9488 - val_loss: 0.3413 - val_accuracy: 0.8997
```

Out[21]: <tensorflow.python.keras.callbacks.History at 0x7f8577e4a4d0>



```
In [33]: 1 model.load_weights('best_weights_epoch215.h5')
2 # model.optimizer = SGD()
3 # keras.backend.set_value(model.optimizer.momentum, 0.7)
4 # keras.backend.set_value(model.optimizer.lr, 0.001)
5
6 path = 'best_weights_epoch225.h5'
7 checkpoint = keras.callbacks.ModelCheckpoint(
8     filepath=path,
9     monitor='val_accuracy',
10    mode='max',
11    save_best_only=True)
12
13 model.compile(loss='categorical_crossentropy', optimizer=SGD(momentum=0.6, learning_rate=0.001),
14               metrics=['accuracy'])
15 model.fit_generator(dagen.flow(X_train, y_train, batch_size),
16                    steps_per_epoch = X_train.shape[0]//batch_size,
17                    epochs = 10, validation_data =(X_test, y_test), callbacks = [checkpoint])
```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/training.py:1940: UserWarning: `Model.fit\_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

warnings.warn("`Model.fit\_generator` is deprecated and "

```
Epoch 1/10
390/390 [=====] - 158s 302ms/step - loss: 0.1525 - accuracy: 0.9464 - val_loss: 0.3418 - val_accuracy: 0.8997
Epoch 2/10
390/390 [=====] - 112s 287ms/step - loss: 0.1480 - accuracy: 0.9478 - val_loss: 0.3422 - val_accuracy: 0.8995
Epoch 3/10
390/390 [=====] - 112s 286ms/step - loss: 0.1490 - accuracy: 0.9479 - val_loss: 0.3419 - val_accuracy: 0.8997
Epoch 4/10
390/390 [=====] - 112s 287ms/step - loss: 0.1480 - accuracy: 0.9484 - val_loss: 0.3413 - val_accuracy: 0.9009
Epoch 5/10
390/390 [=====] - 112s 286ms/step - loss: 0.1488 - accuracy: 0.9481 - val_loss: 0.3405 - val_accuracy: 0.9009
Epoch 6/10
390/390 [=====] - 112s 286ms/step - loss: 0.1487 - accuracy: 0.9476 - val_loss: 0.3414 - val_accuracy: 0.8996
Epoch 7/10
390/390 [=====] - 111s 286ms/step - loss: 0.1461 - accuracy: 0.9488 - val_loss: 0.3421 - val_accuracy: 0.8998
Epoch 8/10
390/390 [=====] - 112s 286ms/step - loss: 0.1444 - accuracy: 0.9498 - val_loss: 0.3410 - val_accuracy: 0.9006
Epoch 9/10
390/390 [=====] - 111s 286ms/step - loss: 0.1465 - accuracy: 0.9492 - val_loss: 0.3383 - val_accuracy: 0.9012
Epoch 10/10
390/390 [=====] - 114s 291ms/step - loss: 0.1496 - accuracy: 0.9467 - val_loss: 0.3405 - val_accuracy: 0.9009
```

Out[33]: <tensorflow.python.keras.callbacks.History at 0x7fbc28df7c50>

```
In [34]: 1 # Test the model
2 score = model.evaluate(X_test, y_test, verbose=1)
3 print('Test loss:', score[0])
4 print('Test accuracy:', score[1])
```

```
313/313 [=====] - 8s 26ms/step - loss: 0.3405 - accuracy: 0.9009
Test loss: 0.340536504983902
Test accuracy: 0.9009000062942505
```

In [ ]: 1

In [ ]:

▶

```
1 # Save the trained weights in to .h5 formath
2 # model.save_weights("DNST_model_100epochs.h6")
3 # print("Saved model to disk")
```

Saved model to disk

After training with 215 epochs we could get val-accuracy of 0.9009

In [ ]:

▶

```
1
```